Learning Discrete Dynamical Systems
under Classification Noise

Zirou Qiu*,[1,2] Zakaria Mehrab*, [1,2] Abhijin Adiga, [2] Madhav V. Marathe,[1,2]

S. S. Ravi,[2,3] Daniel J. Rosenkrantz,[2,3] Richard E. Stearns,[2,3] Anil Vullikanti[1,2] [1]

[1]Computer Science Dept., University of Virginia.
[2]Biocomplexity Institute and Initiative, University of Virginia.
[3]Computer Science Dept., University at Albany – SUNY.

Abstract

Abstract. Due to the important role of discrete dynamical systems in modeling real-world cascading phenomena on networks, problems for learning such systems have garnered considerable attention in ML. However, existing studies on this topic typically assume that the training data is *noise-free*, an assumption that is often impractical. In this work, we address this gap by investigating a more realistic and challenging setting: learning discrete dynamical systems from data contaminated with *noise*. Towards this end, we present *efficient* noise-tolerant learning algorithms that provide provable performance guarantees under the PAC model, and establish tight bounds on sample complexity. We show that, even in the presence of noise, the proposed learner only needs a small training set to infer a system. Notably, the number of training samples required by the algorithm in the noisy setting is the *same* (to within a constant factor) as the information-theoretic upper bound in the noise-free scenario. Further, the number of noisy training samples used by the algorithm is *only a logarithmic factor* higher than the best-known lower bound. Through experimental studies, we evaluate the empirical performance of the algorithms on both synthetic and real-world networks.

Conference version. The paper is **under submission.** Do not distribute.

1 Introduction

Background and Motivation. Discrete dynamical systems serve as formal models for various real-world diffusion processes on networks, including the spread of rumors, information, and diseases [7, 33, 46, 63, 60, 42, 35]. For dynamical systems in the real world, however, one often *cannot* readily obtain a full system specification. Thus, learning unknown components of dynamical systems is an active research area [10, 11, 13, 12, 55, 29, 49, 16, 2].

In essence, a discrete dynamical system consists of an underlying *network* over which a contagion spreads. Vertices in the network are entities such as individuals, and edges denote their relationships. To model a cascade, each vertex has a *contagion state* and an *interaction function*. As the contagion spreads, the states of vertices evolve in discrete time steps based on the mechanism of interaction functions. Thus, the interaction

functions play an important role in the system dynamics as they model the **behavior** of individuals. One concrete example is the *threshold* function, a classic model for social contagions [25, 67, 44, 65, 55, 10]. Here, each individual adopts a contagion (e.g., believes a rumor) only when the number of its neighbors adopting this contagion reaches a tipping point.

Our work focuses on learning discrete dynamical systems where the *interaction functions are unknown*. Notably, existing methods for learning interaction functions (e.g., [2, 53]) were developed under one critical assumption: the training data is *noise-free*. However, it is widely recognized that real-world data are often contaminated by *noisy labels* (i.e., classification noise) [59, 51, 9, 26, 39, 5]. In particular, the presence of noise in training data can significantly degrade the prediction accuracy of learning algorithms that are not noise-tolerant. Nevertheless, noise-tolerant learning for discrete dynamical systems has *not* received attention in the literature. In this work, we address this gap with a systematic study of **learning discrete dynamical systems under classification noise**.

Problem description. There is a target ground-truth discrete dynamical system where all the interaction functions of vertices are **unknown**. A learner must infer the missing functions from only snapshots of the system's dynamics (provided in a training set) which specify the evolution of vertices' states. Moreover, the training set is contaminated by classification noise [5] such that in each snapshot of the dynamics, the updated state of each vertex is incorrect with some (unknown) probability. Our goal is then to design noise-tolerant algorithms that learn all the unknown interaction functions and recover a system that models the behavior of the true system, with performance guarantees under the Probably Approximately Correct (PAC) model [66].

Challenges. Empirical risk minimization (ERM) is often a promising method under the PAC framework. However, the results in [3] show that in our setting, ERM **cannot** be done efficiently unless **P = NP**. In particular, one **cannot** even efficiently approximate the problem of constructing a system that is consistent with the maximum number of training samples. The second difficulty is that, in our noise setting (see Section 2), the probability of a *multi-class label* (a vector of the states of vertices) being *incorrect* in the training set asymptotically **goes to one** as the system size increases. In other words, *almost surely all the multi-class labels in the training set are wrong*.

Given the importance of the problem and the aforementioned challenges, we aim to answer the following questions about learning discrete dynamical systems: 1. *Is efficient learning possible under classification noise*? 2. *How many additional samples do we need compared to the noise-free case*?

Our contributions. Despite the challenges, we answer both questions. In particular, we show that one can still *efficiently* learn discrete dynamical systems under high classification noise. Notably, the number of training samples required is the **same** (to within a constant factor) as the upper bound in the *noise-free* scenario, and it is only a **log factor higher** than the best-known lower bound.

- **Efficient learnability.** Formally, we propose two *efficient* noise-tolerant algorithms V-ERM and VisRange with respective theoretical and empirical advantages. Both algorithms achieve the PAC guarantee: w.p. at least $1 - \delta$, the prediction error is at most ϵ, for any $\epsilon, \delta > 0$. However, Algorithm V-ERM uses $O(n^2 \log(n))$ training samples, whereas Algorithm VisRange uses only $O(n \log n)$ samples.

  From a theoretical perspective, VisRange wins. In particular, the corresponding bound $O(n \log(n))$

on the sample complexity is close to optimal; it matches (to within a constant factor) the information-theoretic *upper* bound in the *noise-free* scenario [27, 41]. That is, the algorithm VisRange is strongly noise-tolerant, with only a marginal increase in the number of samples used in comparison with the noise-free case. Further, its sample complexity is only a factor $O(\log n)$ larger than the general *lower* bound in the noisy setting [5].

- **Experimental evaluation.** We conduct an extensive experimental study of the algorithms over both real and synthetic networks. Our results highlight a drastic difference in their empirical behaviors: VisRange exhibits a *phase transition* w.r.t. the error rate when the size of the training set reaches a critical threshold. This phenomenon is expected from our theoretical analysis. In contrast, V-ERM shows a steady increase in learning accuracy as more training samples are given. Overall, our experiments underscore that V-ERM is the preferred empirical algorithm due to its simplicity and consistent performance. We then further explore the property of V-ERM w.r.t. different model parameters such as network size and error rates.

Related work. The random classification noise (RCN) model is a classic framework for learning under noise, introduced by Angluin and Laird [5] where they proved the efficient learnability of CNF formulas with at most k literals per clause. Many other learning algorithms under the RCN model have been developed for different concept classes (see e.g., [17, 56, 57, 38, 18]). Further, both lower and upper bounds on the sample complexity for learning under classification noise are established in [61, 6, 41, 47]. Several other noise models have also been considered in the literature (e.g., [37, 36, 32, 51, 8, 19]).

Rigorous methods have been proposed for learning various components of a dynamical system, such as the interaction functions, edge parameters, infection sources, and contagion states, from system dynamics [45, 12, 11, 13, 16, 69, 34, 68, 29, 49, 15, 20, 24, 30, 44, 28, 58, 62]. The problem of learning the underlying system topology has also been examined [31, 52, 1, 21, 48, 23, 64]. To the best of our knowledge, the problem of learning discrete dynamical systems under classification noise has *not* been studied in the existing literature.

2 Preliminaries

In this section, we present the formal definitions of discrete dynamical systems over networks and outline our learning setting.

2.1 Discrete dynamical systems

A *discrete dynamical system* over domain $\mathbb{B} = \{0, 1\}$ is defined as a pair $h^* = (\mathcal{G}, \mathcal{F})$, where (i) $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ is an underlying undirected graph with n vertices (e.g., individuals in a social network); (ii) $\mathcal{F} = \{f_v : v \in \mathcal{V}\}$ is a set of *interaction functions*, where f_v is the function for $v \in \mathcal{V}$.

Interaction functions. Each vertex in $\mathcal{G}$ has a *state* from domain $\mathbb{B}$ representing its contagion state (e.g., inactive or active). Starting from the initial states of vertices, a system h^* evolves over *discrete time*, with vertices updating their states *synchronously* using the interaction functions. Specifically, for any interaction function f_v, the inputs are the current states of v's neighbors; the output of f_v is the next state of v. In this

work, we focus on dynamical systems with *threshold* interaction functions. Such systems are fundamental models for the spread of **social contagions** such as rumors and information [25, 67, 44, 65, 55, 10].

Formally, each vertex $v \in \mathcal{V}$ has an integer threshold $\tau_v^* \in [0, \deg_v + 1]$, where $\deg_v$ is the degree of v in $\mathcal{G}$. At each time $t \geq 1$, the function f_v computes v's state at the next time-step $t + 1$ as follows: f_v outputs (state) 1 if the number of state-1 vertices in v's neighborhood at time t is at least τ_v^*; f_v outputs (state) 0 otherwise. In the rumor-spreading example, a person's belief changes when the number of neighbors believing in the rumor reaches a certain tipping point. An example of a threshold dynamical system is shown in Figure 4 in Appendix 7.1.

Configurations. A *configuration* $\mathcal{C}$ of a system h^* is a length-n binary vector specifying the contagion states of all vertices; $\mathcal{C}[v]$ is the state of v under $\mathcal{C}$. Thus, one can view the evolution of system h^* as a time-ordered trajectory of configurations. In such a trajectory, a configuration $\mathcal{C}'$ is the **successor** of $\mathcal{C}$ if the system evolves from $\mathcal{C}$ to $\mathcal{C}'$ in a single time-step; this is denoted by $\mathcal{C}' = h^*(\mathcal{C})$.

2.2 Learning under noise

Let h^* be a ground-truth discrete dynamical system. The learner is provided with *incomplete* information about the system h^*, where the underlying graph is known, but all the interaction functions are **unknown**. In this case, the **hypothesis class** $\mathcal{H}$ consists of all systems with the same graph as h^*, over all possible threshold assignments for vertices in $\mathcal{V}$. By observing the *noisy* snapshots of system dynamics (given in a training set), the learner's goal is to infer the missing interaction functions and learn a system $h \in \mathcal{H}$ that closely approximates the behavior of the true h^*.

The noisy training set. Our algorithms learn from a training set that consists of noisy snapshots of the true system h^*'s dynamics, under the PAC framework. In particular, we extend the well-known Random Classification Noise (RCN) model [5] for binary classification to our multi-class learning context. Let $\mathcal{N} = \{\eta_v : v \in \mathcal{V}(\mathcal{G})\}$ be a collection of **unknown** noise rates where each η_v satisfies $0 < \eta_v \leq \bar{\eta} < 1/2$; here, $\bar{\eta}$ is an upper bound on all noise parameters. Subsequently, a **noisy training set** $\mathcal{T} = \{(\mathcal{C}_j, \hat{\mathcal{C}}_j)\}_{j=1}^q$ is formed as follows. For each data point $(\mathcal{C}_j, \hat{\mathcal{C}}_j) \in \mathcal{T}$,

1. The configuration $\mathcal{C}_j \sim \mathcal{D}$ is sampled independently from an **unknown** distribution $\mathcal{D}$. Let $\mathcal{C}'_j$ be the *true* successor of $\mathcal{C}_j$, produced by the unknown ground truth system h^*.

2. The learner does **not** see the true successor $\mathcal{C}'_j$. Instead, the observed successor $\hat{\mathcal{C}}_j$ in $(\mathcal{C}_j, \hat{\mathcal{C}}_j) \in \mathcal{T}$ is a *noisy* version of $\mathcal{C}'_j$ where the value of each entry $\mathcal{C}'_j[v]$ is *altered* with probability η_v. Formally, for each $v \in \mathcal{V}(\mathcal{G})$: (i) $\hat{\mathcal{C}}_j[v] = \neg\, \mathcal{C}'_j[v]$ w.p. η_v, and (ii) $\hat{\mathcal{C}}_j[v] = \mathcal{C}'_j[v]$ w.p. $1 - \eta_v$.

   As in the RCN model, errors are introduced here by a random process that is independent of the sampling step over $\mathcal{D}$ [5]. Let $\mathcal{O}$ denote the noisy oracle described above. For simplicity, we use $(\mathcal{C}, \hat{\mathcal{C}}) \sim \mathcal{O}$ to denote a training pair generated by the oracle $\mathcal{O}$, such that $\mathcal{C}$ is first sampled from $\mathcal{D}$, then the random noise is applied to $\mathcal{C}'$ to get $\hat{\mathcal{C}}$. We use $\mathcal{T} \sim \mathcal{O}^q$ to denote the sampling of a size-q training set $\mathcal{T}$ from $\mathcal{O}$.

We note that η_v is often fixed w.r.t n in our setting. Thus, in presenting the results on sample complexity in later sections, we sometimes omit the terms involving the noise parameter η_v (or $\bar{\eta}$) for clarity, and mainly

focus on the expressions with the dominating terms given as a function of n.

Remark 1. In our learning under noise model, the learner is **not** involved in generating the training set. Therefore, it does **not** have control over how the noise is applied to each vertex, and it also does **not** know which vertex's state is wrong in any training sample. Further, the learner does **not** have the exact value of each noise parameter η_v, $v \in \mathcal{V}$. It is important to note that our multi-class learning model is different from the standard Random Classification Noise (RCN) model. In particular, in the RCN model, strictly *less than half* of the labels are incorrect in expectation. On the other hand, in our setting, the probability of a multi-class label (i.e., the successor of a configuration) in the training set being *incorrect* (i.e., altered by noise) asymptotically goes to one as n increases. That is, almost surely all the labels in the training set are wrong.

Remark 2. In real-world scenarios, the training set $\mathcal{T}$ can be viewed as a collection of snapshots of the true system dynamics. Consequently, learning from a trajectory of the system evolution can be cast as a special case of our setting. Specifically, suppose $\mathcal{T}$ consists solely of configurations on a particular trajectory P. Then the underlying sampling distribution (unknown to the learner) is such that only configurations on P are sampled with positive probability, while all other configurations are sampled with probability P. In this work, we present learners that work under **arbitrary** sampling distributions, including the one mentioned above.

Learning from the noisy training set. Even though the training set $\mathcal{T}$ is from the noisy oracle $\mathcal{O}$, the aim of the learner *remains to find an appropriate hypothesis w.r.t the **true** (unknown) distribution $\mathcal{D}$.* Given a new $\mathcal{C} \sim \mathcal{D}$, a learned hypothesis h should predict the true successor $\mathcal{C}'$ (i.e., $h^*(\mathcal{C})$) with high probability. Formally, we use $err_\mathcal{D}(h) = \Pr_{\mathcal{C} \sim \mathcal{D}}[h(\mathcal{C}) \neq h^*(\mathcal{C})]$ to denote the **error** admitted by h. Note that in our setting, a prediction is considered incorrect if $h(\mathcal{C})$ and $h^*(\mathcal{C})$ *disagree on the state of at least one vertex*; that is, we want to predict the output states of all vertices correctly.

Following the PAC model, for any parameters $\epsilon, \delta > 0$, a learner should find a hypothesis $h \in \mathcal{H}$ such that with probability at least $1 - \delta$ over $\mathcal{T} \sim \mathcal{O}^q$, the error $err_D(h) \leq \epsilon$. The minimum number of training examples required by any learner to achieve the above PAC-guarantee is known as the **sample complexity** of the class $\mathcal{H}$.

3   Element-wise ERM

We prove the efficient learnability of threshold discrete dynamical systems under classification noise. We begin by establishing the sufficient sample size for a general scheme of multiclass learning under noise using any *element-wise* ERM (defined later). We then propose an **efficient** element-wise ERM algorithm V-ERM for learning threshold discrete dynamical systems that use $O(1/\epsilon^2 \cdot n^2 \cdot \log(n/\delta))$ training samples. Due to space limits, **full proofs appear in the Appendix** 7.2.

A general learning model. We present a general learning model. Following this, learning discrete dynamical systems becomes a *special case* of this general scheme. This new model follows our learning setting in Section 2 for a *finite* **hypothesis class** $\mathcal{H}'$, with the following generalizations: (i) the training set $\mathcal{T}$ consists of q pairs of n-dimensional vectors, denoted by $(\mathcal{W}_j, \hat{\mathcal{W}}_j)$, $j = 1, ..., q$; (ii) the input vector $\mathcal{W}_j \in \{0, 1\}^n$ is drawn from an *unknown* distribution $\mathcal{D}$, where each entry of $\mathcal{W}_j$ is the feature value associated with **an**

entity; (iii) the observed label vector $\hat{\mathcal{W}}_j$ in $(\mathcal{W}_j, \hat{\mathcal{W}}_j)$ is an erroneous version of the true label vector $\mathcal{W}'_j$; (iv) the number of labels for each entity is $k \geq 2$. Also see Appendix 7.2 for a detailed definition of this general learning model.

3.1 Analysis of element-wise ERM for the general model

We analyze the number of training samples required by any element-wise ERM for the general multi-class learning problem defined in the previous section; the result extends the sample complexity proof in [5] for binary classification.

Element-wise ERM. Given a training set $\mathcal{T} = \{(\mathcal{W}_j, \hat{\mathcal{W}}_j)\}_{j=1}^q$, for any hypothesis $h \in \mathcal{H}'$ and entity i, $i = 1, ..., n$, we refer to the number of empirical disagreements $\sum_{j=1}^q \mathbb{1}\left(h(\mathcal{W}_j)[i] \neq \hat{\mathcal{W}}_j[i]\right)$ as the **empirical loss** (over $\mathcal{T}$) of h w.r.t entity i. Let $\mathcal{A}$ be an *element-wise* ERM algorithm. That is, for any training set $\mathcal{T}$, algorithm $\mathcal{A}$ outputs a hypothesis h from the space $\mathcal{H}'$, such that for every entity $i = 1, ..., n$, the empirical loss of h w.r.t i is **minimized**. We note that for some problems, such an algorithm $\mathcal{A}$ might **not** exist. Therefore, the results in this section are for problems that admit at least one such element-wise ERM.

Canonical partition of the hypothesis class. We now to define a partition of the hypothesis class $\mathcal{H}'$. For each ith entity, $i = 1, ..., n$, let $\mathcal{P}_i$ be a partition of the hypothesis class $\mathcal{H}'$ into t_i subsets, denoted by $\mathcal{H}'_1, ..., \mathcal{H}'_{t_i}$, such that the following condition holds for each $\mathcal{H}'_\ell$, $\ell = 1, ..., t_i$: *for any training set $\mathcal{T}$, the empirical loss w.r.t. entity i for all hypotheses in $\mathcal{H}'_\ell$ are the same.*

Note that the construction of such a partition is problem-dependent. Clearly, one trivial partition $\mathcal{P}_i$ is of size $|\mathcal{H}'|$, where each subset in the partition consists of just one hypothesis. Given a collection of partitions $\mathcal{P} = \{\{\mathcal{P}_1, ..., \mathcal{P}_n\}\}$ defined above, let $t_{max}(\mathcal{P})$ be the largest t_i over $i = 1, ..., n$.

Sample complexity. We now establish the sample complexity of learning under the general framework in Lemma 3.1. We note that the main purpose of Lemma 3.1 is to later derive a sample complexity bound for learning discrete dynamical systems. Nevertheless, this general result may also be of independent interest.

> **Lemma 3.1.** *Let $\mathcal{P} = \{\{\mathcal{P}_1, ..., \mathcal{P}_n\}\}$ be a collection of canonical partitions of the hypothesis class $\mathcal{H}'$; $t_{max}(\mathcal{P})$ is the size of the largest partition in $\mathcal{P}$. For any $\epsilon, \delta \in (0, 1)$, and any $\eta_i \leq \bar{\eta} < 1/2$, $i = 1, ..., n$, with a training set of size*
>
> $$q = O\left(\frac{\bar{\eta}}{(1 - 2\bar{\eta})^2} \cdot \frac{1}{\epsilon^2} \cdot n^2 \log\left(\frac{t_{max}(\mathcal{P}) \cdot n}{\delta}\right)\right)$$
>
> *any element-wise ERM learns a $h \in \mathcal{H}'$ such that with probability at least $1 - \delta$ (over $\mathcal{T} \sim \mathcal{O}^q$), $err_{\mathcal{D}}(h) < \epsilon$.*

Remark 3. We note the following: *the element-wise ERM $\mathcal{A}$ in general does **not** minimize the empirical loss over the training set $\mathcal{T}$*. In fact, due to the high probability of receiving a wrong prediction on each sampled data, any true ERM is unlikely to perform well in this case. The second remark is that knowing the exact value of $\bar{\eta}$ is **not** needed in our setting. In particular, one can easily extend the binary search technique introduced in [5] (Theorem 3 in [5]) to estimate the value of $\bar{\eta}$, such that w.p. at least $1 - \delta$, the estimated $\bar{\eta}$ satisfies $\eta_i \leq \bar{\eta} < 1/2$ for all $i = 1, ..., n$.

3.2 An efficient vertex-wise ERM for learning discrete dynamical systems

Given Lemma 3.1, what remains is to (i) find an actual (efficient) algorithm that is an element-wise ERM and (ii) determine an appropriate partition of the hypothesis class $\mathcal{H}$ for discrete dynamical systems. In this section, we answer these two questions for learning discrete dynamical systems. In particular, we present a simple and efficient algorithm that learns a hypothesis (system) that minimizes the empirical risk w.r.t *each vertex* over the training set $\mathcal{T}$. Further, we show the existence of a canonical partition of $\mathcal{H}$ w.r.t each vertex $v \in \mathcal{V}$, such that the size of each $\mathcal{P}_v$ is at most $\Delta + 2$, where Δ is the maximum degree of the graph. Subsequently, we prove that the number of samples needed by the algorithm is $O(1/\epsilon^2 \cdot n^2 \cdot \log(n/\delta))$.

The algorithm V-ERM. A system h is inferred as follows: for each vertex $v \in \mathcal{V}(\mathcal{G})$, we learn τ_v to be the value such that the number of disagreements $\sum_{j=1}^{q} \mathbb{1}\left(h(\mathcal{C}_j)[v] \neq \hat{\mathcal{C}}_j[v]\right)$ over the training set $\mathcal{T}$ is **minimized**. Note that such a τ_v can be found in polynomial time by iterating over each integer in $[0, \deg_v + 1]$. The pseudocode appears as Algorithm 1 in Appendix 7.3.

The canonical partition. For each vertex v, a canonical partition $\mathcal{P}_v = \{\mathcal{H}_0, ..., \mathcal{H}_{\deg_v+1}\}$ of $\mathcal{H}$ is constructed such that each hypothesis in $\mathcal{H}_\ell$, $\ell = 0, ..., \deg_v + 1$, has the following property: *the threshold of v is ℓ*. One can verify that, for any training set, the empirical loss w.r.t. vertex v for all hypotheses in $\mathcal{H}_\ell$ are the same. Further, the size of the partition is $\deg_v + 2$. Consequently, $t_{\max}(\mathcal{P}) = \Delta + 2 \leq n + 1$, where $\mathcal{P} = \{\{\mathcal{P}_1, ..., \mathcal{P}_n\}\}$.

Lastly, for our problem of learning threshold discrete dynamical systems, the label for each vertex is either 0 or 1; that is, $k = 2$. By Lemma 3.1, the Theorem follows.

> **Theorem 3.2.** *For any $\epsilon, \delta \in (0, 1)$, and any $\eta_v \leq \bar{\eta} < 1/2, v \in \mathcal{V}$, with a training set $\mathcal{T}$ of size*
>
> $$q = O\left(\frac{\bar{\eta}}{(1 - 2\bar{\eta})^2} \cdot \frac{1}{\epsilon^2} \cdot n^2 \cdot \log(\frac{n}{\delta})\right)$$
>
> *Algorithm 1 (i.e., V-ERM) learns a hypothesis $h \in \mathcal{H}$ such that with probability at least $1 - \delta$ (over $\mathcal{T} \sim \mathcal{O}^q$), $err_{\mathcal{D}}(h) < \epsilon$.*

Remark 4. Theorem 3.2 establishes the efficient learnability of threshold discrete dynamical systems under classification noise. Note that the sample complexity bound derived from V-ERM is not optimal, as it follows from the more general result, namely Lemma 3.1. In the next section, we present a more sophisticated algorithm with much lower sample complexity. Nevertheless, it is noteworthy that such a simple algorithm (i.e., V-ERM) can already achieve PAC performance guarantees using only $O(n^2 \log(n))$ noisy samples. The simplicity of this algorithm makes it well-suited for use in practice. In particular, one key property of V-ERM discovered through empirical analysis is that, *the learning accuracy consistently increases as more training samples are provided*. Such a property is *not* observed for the more complex algorithms presented in the next section.

4 Learning based on Visiting Times

One immediate question is whether the sample complexity established in the previous section can be improved. In this section, we answer the question with a more problem-dependent analysis and improve the bound so that it **matches the general information-theoretic upper bound in noise-free scenarios**. Toward this end, we present the algorithm VisRange that uses $O(1/\epsilon \cdot n \log(n/\delta))$ training samples under our high-noise setting, which is only a constant factor larger than the general upper bound for learning under the noise-free scenario [27]. That is, despite the presence of noise, the algorithm achieves efficient learnability without needing many more samples compared to the noise-free scenario. Further, *our established upper bound is only a factor $O(\log n)$ larger than the best-known lower bound* [6]. Due to space limit, **all proofs appear in Appendix 7.3**.

We first discuss a simplified version of the algorithm for ease of understanding. This simpler algorithm uses $O(1/\epsilon \cdot \Delta n \log(n/\delta))$ training samples, where Δ is the maximum degree of the graph.

Visiting a score. We define the notion of *visiting a score*, which plays a critical role in the algorithm. Let v be a vertex in the graph. Given a configuration $\mathcal{C}$, the **score** of v under $\mathcal{C}$, denoted by $\texttt{score}(\mathcal{C}, v)$, is the number of state-1 neighbors of v in $\mathcal{C}$. Note that the score of v is in the range $[0, \deg_v + 1]$. For a $\mathcal{C} \sim \mathcal{D}$, we say that a score $s \in [0, \deg_v + 1]$ is **visited** by $\mathcal{C}$ for v if $\texttt{score}(\mathcal{C}, v) = s$. Subsequently, the **visiting probability** of a score s w.r.t. v is the probability of visiting s over $\mathcal{C} \sim \mathcal{D}$. Given a training set $\mathcal{T} \sim \mathcal{O}^q$, the **visiting time** of a score s is a random variable recording the number of times s got visited by $\mathcal{C}$, summing over all pairs $(\mathcal{C}, \hat{\mathcal{C}}) \in \mathcal{T}$.

The simplified algorithm VisScore. Fix a vertex $v \in \mathcal{V}$. At a high level, when the size q of the training set $\mathcal{T}$ is sufficiently large, with probability at least $1 - \delta$, each score with a high visiting probability will be visited a sufficiently large number of times in the training set. We then learn from the majority vote over the output states for each such input score in $\mathcal{T}$.

Formally, let S be the set of scores that are visited at least $q \cdot \epsilon/(2\Delta n)$ times in $\mathcal{T}$. For each score $s \in S$, the algorithm computes the *majority* output state (break ties randomly) of v over the erroneous successors under input score s; let ℓ'_s be the majority output state for $s \in S$ over $\mathcal{T}$. Lastly, the threshold of v is learned as $\tau_v = 1 + \max\{s \in S : \ell'_s = 0\}$; $\tau_v = 0$ if $\ell'_s = 1$ for all $s \in S$. This procedure is then performed for each vertex $v \in \mathcal{V}$, which produces a system with all thresholds being inferred. The pseudocode for the algorithm is shown in Algorithm 2 in Appendix 7.3.

In Theorem 4.1, we show that Algorithm 2 (i.e., VisScore) uses at most $O(\Delta n \log(n))$ samples to infer all the thresholds. For graphs where the maximum degree is $o(n)$, this bound is already better than the $O(n^2 \log(n))$ bound for the element-wise ERM Algorithm 1 (i.e., V-ERM) in Section 3.

Theorem 4.1. *For any $\epsilon, \delta \in (0, 1)$, and any $\eta_v \leq \bar{\eta} < 1/2, v \in \mathcal{V}$, with a training set of size*

$$q = O\left(\frac{1 - \bar{\eta}}{(1/2 - \bar{\eta})^2} \cdot \frac{1}{\epsilon} \cdot \Delta n \cdot \log(\frac{n}{\delta})\right)$$

Algorithm 2 (i.e., VisScore) learns a hypothesis $h \in \mathcal{H}$ such that with probability at least $1 - \delta$ (over $\mathcal{T} \sim \mathcal{O}^q$), we have $err_{\mathcal{D}}(h) < \epsilon$.

4.1 The algorithm based on visiting times of ranges

In this section, we extend the simplified algorithm (i.e., VisScore) to our final algorithm VisRange which uses $O(1/\epsilon \cdot n \log{(n/\delta)})$ training samples. We then show that this bound is tight by comparing it with the general upper and lower bounds from the literature.

Visiting a range. We extend the notion of visiting a score (used in VisScore) to *visiting a range of scores*. Let v be a vertex in the graph. Let $R_{s_1,s_2} = [s_1, s_2]$ be a range of scores for v, $s_1, s_2 \in \{0, ..., \deg_v + 1\}$, $s_1 \leq s_2$. Note that there are $O(\Delta^2)$ such ranges for each v. For a configuration $\mathcal{C} \sim \mathcal{D}$, we say that a range R_{s_1,s_2} is **visited** by $\mathcal{C}$ for v if $\texttt{score}(\mathcal{C}, v) \in R_{s_1,s_2}$. Similarly, the **visiting probability** of a range R_{s_1,s_2} w.r.t. v is the probability of visiting R_{s_1,s_2} over $\mathcal{C} \sim \mathcal{D}$. Lastly, the **visiting time** of a range R_{s_1,s_2} is a random variable representing the number of times R_{s_1,s_2} got visited by $\mathcal{C}$, summing over all pairs $(\mathcal{C}, \hat{\mathcal{C}}) \in \mathcal{T}$.

The algorithm VisRange**.** The algorithm is an extended version of VisScore, where we now care about ranges of scores being visited. Fix a vertex $v \in \mathcal{V}$. Intuitively, when the size q of the training set $\mathcal{T}$ is sufficiently large, with probability at least $1 - \delta$, each range of scores (for v) with a high visiting probability would be visited a large number of times. We then learn from the majority vote of the output states of v over all the input scores in each such range with high visiting probabilities.

Let S be the set of ranges that are visited at least $\epsilon/(2n) \cdot q$ times in $\mathcal{T}$. For each range $(s_1, s_2) \in S$, the algorithm counts the total number of output state-0 and output state-1 over all the erroneous successors with input scores in $[s_1, s_2]$; let ℓ'_{s_1,s_2} be the corresponding majority output state for the range $[s_1, s_2]$ over $\mathcal{T}$. Lastly, the threshold of v is learned as $\tau_v = 1 + \max\{s_1 : (s_1, s_2) \in S \text{ and } \ell'_{s_1,s_2} = 0\}$. If $\ell'_{s_1,s_2} = 1$ for all $(s_1, s_2) \in S$, then $\tau_v = 0$. The pseudocode of the algorithm is shown in Algorithm 3 in Appendix 7.3. It is clear that the overall algorithm runs in polynomial time. In Theorem 4.2, we show that $O(n \log{(n)})$ samples are sufficient for Algorithm 3 to learn the system.

Theorem 4.2. *For any $\epsilon, \delta \in (0, 1)$, and any $\eta_v \leq \bar{\eta} < 1/2$, $v \in \mathcal{V}$, with a training set of size*

$$q = O\left(\frac{1 - \bar{\eta}}{(1/2 - \bar{\eta})^2} \cdot \frac{1}{\epsilon} \cdot n \cdot \log(\frac{n}{\delta})\right)$$

Algorithm 3 (i.e., VisRange) learns a hypothesis $h \in \mathcal{H}$ such that with probability at least $1 - \delta$ (over $\mathcal{T} \sim \mathcal{O}^q$), we have $err_{\mathcal{D}}(h) < \epsilon$.

4.2 Tightness of the sample complexity bound

We now discuss the tightness of our bound $O\left(\frac{1-\bar{\eta}}{(1/2-\bar{\eta})^2} \cdot \frac{1}{\epsilon} \cdot n \cdot \log(\frac{n}{\delta})\right)$ established in Theorem 4.2 by comparing it with the general upper and lower bounds on PAC learning.

- **Upper bound.** The work by Haussler [27] shows that the sample complexity of general PAC learning

in the **noise-free** setting is $O\left(\frac{1}{\epsilon} \cdot \log\left(\frac{|\mathcal{H}|}{\delta}\right)\right)$, where $|\mathcal{H}|$ is the size of the hypothesis class. For learning threshold dynamical systems, since $|\mathcal{H}| = O(n^n)$, the corresponding **noise-free** upper bound becomes $O\left(\frac{1}{\epsilon} \cdot n \log\left(\frac{n}{\delta}\right)\right)$.

In the presence of noise, one would expect the above bound to become higher. Nevertheless, our derived bound under classification noise, $O\left(\frac{1-\bar{\eta}}{(1/2-\bar{\eta})^2} \cdot \frac{1}{\epsilon} \cdot n \cdot \log(\frac{n}{\delta})\right)$, is **only a factor** $O(1/(1-\bar{\eta})^2)$ larger than the noise-free upper bound, where the expressions involving the dominant term n remain the same. Our result also matches a general upper bound of $O\left(\frac{1}{(1-2\eta)^2} \cdot \frac{1}{\epsilon} \cdot n \log\left(\frac{n}{\delta}\right)\right)$ by Laird [41] for PAC learning under classification noise.

More importantly, **one cannot generally expect efficient PAC learning algorithms to achieve the upper bounds** from Haussler [27] or Laird [41] since there exist problems that are not efficiently PAC learnable unless **NP = RP** [39]. Our result (i.e., Theorem 4.2), using VisRange reveals that such a bound indeed holds for efficiently learning threshold dynamical systems.

- **Lower bound.** By Aslam and Decatur [6], a *general* lower bound on sample complexity for PAC learning under classification noise is $\Omega\left(\frac{1}{(1-2\eta)^2} \cdot \frac{1}{\epsilon} \cdot \left(\text{Ndim}(\mathcal{H}) + \log\frac{1}{\delta}\right)\right)$, where $\text{Ndim}(\mathcal{H})$ is the Natarajan dimension [50] of the class $\mathcal{H}$. For learning threshold systems, it is known that $\text{Ndim}(\mathcal{H}) = n$ [54]. Hence, the lower bound becomes $\Omega\left(\frac{1}{(1-2\eta)^2} \cdot \frac{1}{\epsilon} \cdot \left(n + \log\frac{1}{\delta}\right)\right)$. Notably, our upper bound in Theorem 4.2 is **only a factor** $O(\log n)$ higher than the general lower bound.

Remark 5. With everything in place, we now discuss the advantages of V-ERM (i.e., Algorithm 1) and VisRange (i.e., Algorithm 3) from both theory and practice perspectives:

- **Theory.** With more problem-dependent mechanism and analysis as given for VisRange, one can significantly improve the upper bound (by a multiplicative factor of n) of the sufficient training size from the bound derived by V-ERM. Specifically, there exists a problem instance (e.g., a graph $\mathcal{G}$ and a distribution $\mathcal{D}$) such that VisRange requires multiplicative $\Omega(n)$ *fewer* samples than V-ERM to achieve the same error rate. Further, VisRange provides an important (and somewhat surprising) theory insight: despite the presence of classification noise, one can still *efficiently* learn threshold dynamical systems using the *same* number of training samples (to within only a constant factor) as the noise-free case. Further, the number of required samples is only a factor $O(\log n)$ higher than the lower bound.

- **Practice.** A distinct nature of VisRange is that, one expects to see a drop in the error rate only when the size of the training set q reaches a critical threshold (i.e., a phase transition). Before q reaches that threshold, however, the behavior of the algorithm can be unpredictable as the high-probability guarantee is not yet satisfied. On the other hand, such a critical threshold on the training set size is **not** inherent to V-ERM. In particular, as more training samples are given, the corresponding loss for V-ERM always decreases correspondingly. Thus, V-ERM is the preferred method in practice due to both its simplicity and smoother performance. As a result, we focus on V-ERM for our empirical evaluations, as shown in the next section.

5 Experimental Evaluation

In this section, we present experimental studies on the practical feasibility of the proposed algorithms, V-ERM and VisRange, on both synthetic and real-world networks [22, 43, 40]. Our goal is to complement our theoretical results by analyzing the empirical behaviors of our algorithms across different model parameters.

Experimental setup. The details of the synthetic networks (`Gnp`) and real-world networks are given in Appendix 7.4. For each network, we create a target ground-truth system h^* where the thresholds are unknown to the learning algorithm. In particular, the threshold τ_v for each vertex v in h^* was chosen uniformly randomly over the range $[0, \deg(v) + 1]$. Under each h^*, a training set $\mathcal{T} = \{(\mathcal{C}_j, \hat{\mathcal{C}}_j)\}_{j=1}^q$ is generated such that (i) in each $\mathcal{C}_j$, the state of a vertex is 0 or 1 with the same probability; (ii) $\hat{\mathcal{C}}_j$ is a noisy version of the true successor $\mathcal{C}'_j$ where the noise rate $\eta > 0$ is the same for all vertices. Here, we consider different η values ranging from 0.05 to 0.4. Given a training set $\mathcal{T}$, our algorithm then learns a system h by inferring all the thresholds. Lastly, to quantify the solution quality, we sample $1,000$ new configurations and compute the **empirical loss** ℓ, which is the proportion of the sampled configurations that the learned hypothesis h makes incorrect predictions.

The parameter settings are given in Table 1 in the Appendix. Overall, we generated up to $5,000$ training samples for learning each system under synthetic networks with up to 4000 vertices and varying density. For other networks, the number of training samples is up to twice the network size.

5.1 Experimental Results

Figure 1: Performance comparison of V-ERM and VisRange across various network sizes. The average loss is shown after every 50 training samples. Across all the experiments, noise η is set to 0.05 and average density d_{avg} is set to 5.

Distinct behaviors of the algorithms. We first examine the behaviors of the algorithms in terms of the loss ℓ when more training samples are given. Figure 1 shows the change of loss for V-ERM (i.e., Algorithm 1) and VisRange (i.e., Algorithm 3) for synthetic networks of different sizes n. We observe that for larger n, VisRange exhibits a phase transition w.r.t. loss when the number of training samples q reaches a critical threshold. Algorithm VisScore (i.e., Algorithm 2) also exhibits a similar phase transition behavior (see Figure 6 in the Appendix). Such phenomena are expected from our theoretical analysis (see Remark 5). In contrast, V-ERM steadily gains in performance when it is given more training samples.

From Figure 1, we further observe that VisRange achieves a loss lower than V-ERM with fewer training samples in the long run for network sizes of 2000 or 4000. Nonetheless, as the network size grows, the required number of training samples for VisRange to reach an acceptable empirical loss also grows. Overall, V-ERM is more suitable for practical purposes, where there is often a trade-off between performance and the number of samples. Since the learning performance of V-ERM steadily improves with increase in the

number of samples, such a trade-off can be obtained with fewer training samples with V-ERM compared to VisRange. Overall, although VisRange can promise an improved performance in the long run, V-ERM provides reasonable performance in most practical scenarios. Therefore, the remaining experiments focus on V-ERM.

(a) Synthetic Networks   (b) Real-world Networks

Figure 2: Learning curve of V-ERM (Algorithm 1) for various graphs under different noise settings. The shaded region accounts for one standard deviation error.

Further evaluation of V-ERM**.** We perform additional studies on the more applicable algorithm V-ERM. The first evaluation highlights its learning behavior across different noise settings. The results are shown in Figure 2 in two parts, for synthetic networks and real-world networks, respectively.

Figure 2a shows the loss of V-ERM for two synthetic networks under different noise settings. Overall, when the noise rate η increases, we see an increase in the number of training samples needed by V-ERM to achieve the same loss. This behavior is expected. The first plot in Figure 2a is for a sparse network with 500 vertices. We find that even under the high noise setting ($\eta = 0.4$), the algorithm performs well. For instance, when $\eta \leq 0.3$, the loss is close to 0 after observing fewer than $3,000$ training samples. The second plot in Figure 2a is for a network with $5,000$ vertices and an average degree of 40, a larger and much denser network. Even in such case, V-ERM achieves reasonable accuracy for moderate noise levels ($\eta \leq 0.2$). Experimental results for a more comprehensive set of synthetic networks are presented in Figure 8 in Appendix 7.4. Similar results for the two real-world networks are shown in Figure 2b. Overall, we observe that V-ERM achieves a high accuracy upon observing samples that are much less than its worst-case theoretical bound (shown in Theorem 3.2).

(a) With varying noise   (b) With varying graph size   (c) With varying density
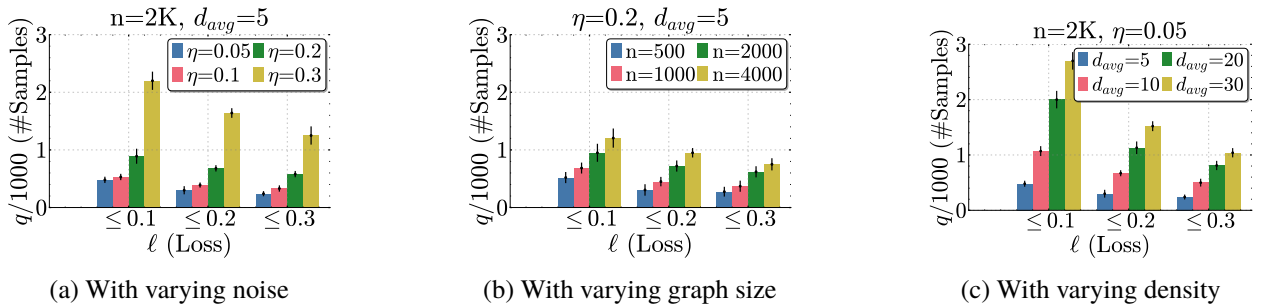
Figure 3: Number of training samples needed for V-ERM for achieving a specified loss. Error bars account for two standard deviation errors.

From the observed decreasing pattern of the loss curve, one can conclude that V-ERM will converge towards a zero-loss as it is provided with more training samples. We illustrate this fact using an alternate perspective in Figure 3, where we show the number of training samples required to achieve loss below a certain threshold under three settings: (i) various noise rates (Figure 3a), (ii) various network sizes (Figure 3b), and

(iii) various graph densities (Figure 3c). We observe that as the threshold of acceptable loss decreases, the number of required training samples increases. Interestingly, this phenomenon is most sensitive to density, followed by noise and then the network size. Results from a similar experiment on real-world networks are presented in Figure 7 in Appendix 7.4.

6 Future Work

In this work, we presented efficient algorithms for learning threshold dynamical systems under classification noise. Much work remains to be done for learning discrete dynamical systems under different models of noise. First, it is of interest to investigate whether our algorithms can be extended to other classes of interaction functions such as weighted threshold functions and symmetric functions [14]. A second direction is to investigate whether the sample complexity bounds can be improved when additional information about a dynamical system (e.g., correct threshold values of some vertices) is available. Finally, it is also of interest to study whether improved learning algorithms can be developed for restricted graph topologies such as planar graphs or intersection graphs.

References

[1] B. Abrahao, F. Chierichetti, R. Kleinberg, and A. Panconesi. Trace complexity of network inference. In *Proceedings of the 19th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 491–499. ACM, 2013.

[2] Abhijin Adiga, Chris J Kuhlman, Madhav Marathe, S Ravi, and Anil Vullikanti. PAC learnability of node functions in networked dynamical systems. In *International Conference on Machine Learning*, pages 82–91. PMLR, 2019.

[3] Abhijin Adiga, Chris J Kuhlman, Madhav V Marathe, S S Ravi, Daniel J Rosenkrantz, and Richard E Stearns. Inferring local transition functions of discrete dynamical systems from observations of system behavior. *Theoretical Computer Science*, 679:126–144, 2017.

[4] N. Ahmed, R. Alo, C. Amelink, Y. Y. Baek, A. Chudhary, K. Collinsand A. Esterline, E. Fox, G. Fox, A. Hagberg, R. Kenyon, C. J. Kuhlman, J. Leskovec, D. Machi, M. V. Marathe, M. Meghanathan, Y. Miyazaki, J. Qiu, N. Ramakrishnan, S. S. Ravi, R. Rossi, R. Sosic, and Gregor von Laszewski. `net.science`: A cyberinfrastructure for sustained innovation in network science and engineering. In *Proc. Gateways Conference*, pages 1–4, Oct. 2020.

[5] Dana Angluin and Philip Laird. Learning from noisy examples. *Machine learning*, 2:343–370, 1988.

[6] Javed A. Aslam and Scott E. Decatur. On the sample complexity of noise-tolerant learning. *Information Processing Letters*, 57(4):189–195, 1996.

[7] Federico Battiston, Giulia Cencetti, Iacopo Iacopini, Vito Latora, Maxime Lucas, Alice Patania, Jean-Gabriel Young, and Giovanni Petri. Networks beyond pairwise interactions: structure and dynamics. *Physics Reports*, 874:1–92, 2020.

[8] Nader H. Bshouty, Nadav Eiron, and Eyal Kushilevitz. PAC learning with nasty noise. *Theoretical Computer Science*, 288:255–275, 2002.

[9] Nicolo Cesa-Bianchi, Eli Dichterman, Paul Fischer, Eli Shamir, and Hans Ulrich Simon. Sample-efficient strategies for learning in the presence of noise. *Journal of the ACM (JACM)*, 46(5):684–719, 1999.

[10] Wei Chen, Xiaoming Sun, Jialin Zhang, and Zhijie Zhang. Network inference and influence maximization from samples. In *International Conference on Machine Learning*, pages 1707–1716. PMLR, 2021.

[11] Yanxi Chen and H. Vincent Poor. Learning mixtures of linear dynamical systems. *International Conference on Machine Learning*, 162:3507–3557, 2022.

[12] Vincent Conitzer, Debmalya Panigrahi, and Hanrui Zhang. Learning opinions in social networks. In *International Conference on Machine Learning*, pages 2122–2132. PMLR, 2020.

[13] Vincent Conitzer, Debmalya Panigrahi, and Hanrui Zhang. Learning influence adoption in heterogeneous networks. In *Proc. AAAI*, pages 6411–6419. AAAI, 2022.

[14] Y. Crama and P. L. Hammer. *Boolean Functions: Theory, Algorithms and Applications*. Cambridge University Press, New York, NY, 2011.

[15] Hadi Daneshmand, Manuel Gomez-Rodriguez, Le Song, and Bernhard Schoelkopf. Estimating diffusion network structures: Recovery conditions, sample complexity & soft-thresholding algorithm. In *International conference on machine learning*, pages 793–801. PMLR, 2014.

[16] Quinlan E Dawkins, Tianxi Li, and Haifeng Xu. Diffusion source identification on networks with statistical confidence. In *International Conference on Machine Learning*, pages 2500–2509. PMLR, 2021.

[17] Scott E. Decatur. PAC learning with constant-partition classification noise and applications to decision tree induction. In Douglas H. Fisher, editor, *Proceedings of the Fourteenth International Conference on Machine Learning (ICML 1997), Nashville, Tennessee, USA, July 8-12, 1997*, pages 83–91. Morgan Kaufmann, 1997.

[18] Scott E. Decatur and Rosario Gennaro. On learning from noisy and incomplete examples. In Wolfgang Maass, editor, *Proceedings of the Eigth Annual Conference on Computational Learning Theory, COLT 1995, Santa Cruz, California, USA, July 5-8, 1995*, pages 353–360. ACM, 1995.

[19] Ilias Diakonikolas, Themis Gouleakis, and Christos Tzamos. Distribution-independent PAC learning of halfspaces with Massart noise. In *Advances in Neural Information Processing Systems*, volume 32, pages 4751–4762, 2019.

[20] Nan Du, Yingyu Liang, Maria Balcan, and Le Song. Influence function learning in information diffusion networks. In *International Conference on Machine Learning*, pages 2016–2024. PMLR, 2014.

[21] Nan Du, Le Song, Ming Yuan, and Alex Smola. Learning networks of heterogeneous influence. *Advances in neural information processing systems*, 25, 2012.

[22] P. Erdös and A. Renyi. On random graphs I. *Publ. math. debrecen*, 6(18):290–297, 1959.

[23] M. Gomez-Rodriguez, J. Leskovec, and A. Krause. Inferring networks of diffusion and influence. In *Proceedings of the 16th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1019–1028. ACM, 2010.

[24] S. González-Bailón, J. Borge-Holthoefer, A. Rivero, and Y. Moreno. The dynamics of protest recruitment through an online network. *Scientific Reports*, 1:7 pages, 2011.

[25] Mark Granovetter. Threshold models of collective behavior. *American journal of sociology*, 83(6):1420–1443, 1978.

[26] Shivani Gupta and Atul Gupta. Dealing with noise problem in machine learning data-sets: A systematic review. *Procedia Computer Science*, 161:466–474, 2019.

[27] David Haussler. Quantifying inductive bias: AI learning algorithms and Valiant's learning framework. *Artificial intelligence*, 36(2):177–221, 1988.

[28] Shushan He, Hongyuan Zha, and Xiaojing Ye. Network diffusions via neural mean-field dynamics. *Advances in Neural Information Processing Systems*, 33:2171–2183, 2020.

[29] Xinran He, Ke Xu, David Kempe, and Yan Liu. Learning influence functions from incomplete observations. *Advances in Neural Information Processing Systems*, 29, 2016.

[30] Lisa Hellerstein and Rocco A Servedio. On PAC learning algorithms for rich boolean function classes. *Theoretical Computer Science*, 384(1):66–76, 2007.

[31] Hao Huang, Qian Yan, Lu Chen, Yunjun Gao, and Christian S Jensen. Statistical inference of diffusion networks. *IEEE Transactions on Knowledge and Data Engineering*, 33(2):742–753, 2019.

[32] Shahin Jabbari, Robert C. Holte, and Sandra Zilles. PAC-learning with general class noise models. In Birte Glimm and Antonio Krüger, editors, *KI 2012: Advances in Artificial Intelligence - 35th Annual German Conference on AI, Saarbrücken, Germany, September 24-27, 2012. Proceedings*, volume 7526 of *Lecture Notes in Computer Science*, pages 73–84. Springer, 2012.

[33] Zhiwei Ji, Ke Yan, Wenyang Li, Haigen Hu, and Xiaoliang Zhu. Mathematical and computational modeling in complex biological systems. *BioMed research international*, 2017, 2017.

[34] Dimitris Kalimeris, Yaron Singer, Karthik Subbian, and Udi Weinsberg. Learning diffusion using hyperparameters. In *International Conference on Machine Learning*, pages 2420–2428. PMLR, 2018.

[35] S. Kauffman, C. Peterson, B. Samuelsson, and C. Troein. Random Boolean network models and the yeast transcriptional network. *Proc. National Academy of Sciences (PNAS)*, 100(25):14796–14799, Dec. 2003.

[36] Michael J. Kearns. Efficient noise-tolerant learning from statistical queries. In S. Rao Kosaraju, David S. Johnson, and Alok Aggarwal, editors, *Proceedings of the Twenty-Fifth Annual ACM Symposium on Theory of Computing, May 16-18, 1993, San Diego, CA, USA*, pages 392–401. ACM, 1993.

[37] Michael J. Kearns and Ming Li. Learning in the presence of malicious errors. *SIAM J. Comput.*, 22(4):807–837, 1993.

[38] Michael J. Kearns and Robert E. Schapire. Efficient distribution-free learning of probabilistic concepts (extended abstract). In *31st Annual Symposium on Foundations of Computer Science, St. Louis, Missouri, USA, October 22-24, 1990, Volume I*, pages 382–391. IEEE Computer Society, 1990.

[39] Michael J Kearns and Umesh Vazirani. *An introduction to computational learning theory*. MIT press, 1994.

[40] Jérôme Kunegis. Konect: the Koblenz network collection. In *Proceedings of the 22nd international conference on world wide web*, pages 1343–1350, 2013.

[41] Philip D Laird. *Learning from good and bad data*, volume 47. Springer Science & Business Media, 2012.

[42] R. Laubenbacher and B. Stigler. A computational algebra approach to the reverse engineering of gene regulatory networks. *J. Theoretical Biology*, 229:523–537, 2004.

[43] Jure Leskovec, Jon Kleinberg, and Christos Faloutsos. Graph evolution: Densification and shrinking diameters. *ACM transactions on Knowledge Discovery from Data (TKDD)*, 1(1):2–es, 2007.

[44] Shuai Li, Fang Kong, Kejie Tang, Qizhi Li, and Wei Chen. Online influence maximization under linear threshold model. *Advances in Neural Information Processing Systems*, 33:1192–1204, 2020.

[45] Andrey Lokhov. Reconstructing parameters of spreading models from partial observations. *Advances in Neural Information Processing Systems*, 29, 2016.

[46] Kristian Lum, Samarth Swarup, Stephen Eubank, and James Hawdon. The contagious nature of imprisonment: an agent-based model to explain racial disparities in incarceration rates. *Journal of The Royal Society Interface*, 11(98):2014.0409, 2014.

[47] Sumona Mukhopadhyay and Santo Banerjee. Learning dynamical systems in noise using convolutional neural networks. *Chaos: An Interdisciplinary Journal of Nonlinear Science*, 30(10), 2020.

[48] Seth Myers and Jure Leskovec. On the convexity of latent social network inference. *Advances in neural information processing systems*, 23, 2010.

[49] H. Narasimhan, D. C. Parkes, and Y. Singer. Learnability of influence in networks. In *Advances in Neural Information Processing Systems*, pages 3186–3194, 2015.

[50] Balas K Natarajan. On learning sets and functions. *Machine Learning*, 4(1):67–97, 1989.

[51] Nagarajan Natarajan, Inderjit S Dhillon, Pradeep K Ravikumar, and Ambuj Tewari. Learning with noisy labels. *Advances in neural information processing systems*, 26, 2013.

[52] Jean Pouget-Abadie and Thibaut Horel. Inferring graphs from cascades: A sparse recovery framework. In *International Conference on Machine Learning*, pages 977–986. PMLR, 2015.

[53] Z. Qiu, A. Adiga, M. V. Marathe, S. S. Ravi, D. J. Rosenkrantz, R. E. Stearns, and A. Vullikanti. Learning the topology and behavior of discrete dynamical systems. In *Proc. 38th AAAI Conference on Artificial Intelligence (AAAI-2024)*, pages 14722–14730, 2024.

[54] Zirou Qiu, Abhijin Adiga, Madhav V Marathe, S S Ravi, Daniel J Rosenkrantz, Richard E Stearns, and Anil Vullikanti. Efficient PAC learnability of dynamical systems over multilayer networks. ArXiv: https://arxiv.org/abs/2405.06884, 2024.

[55] Daniel J Rosenkrantz, Abhijin Adiga, Madhav Marathe, Zirou Qiu, SS Ravi, Richard Stearns, and Anil Vullikanti. Efficiently learning the topology and behavior of a networked dynamical system via active queries. In *International Conference on Machine Learning*, pages 18796–18808. PMLR, 2022.

[56] Yasubumi Sakakibara. On learning from queries and counterexamples in the presence of noise. *Inf. Process. Lett.*, 37(5):279–284, 1991.

[57] Yasubumi Sakakibara. Noise-tolerant occam algorithms and their applications to learning decision trees. *Mach. Learn.*, 11:37–62, 1993.

[58] Augusto Santos, Diogo Rente, Rui Seabra, and José MF Moura. Learning the causal structure of networked dynamical systems under latent nodes and structured noise. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 14866–14874, 2024.

[59] Fahad Sarfraz, Elahe Arani, and Bahram Zonooz. Noisy concurrent training for efficient learning under label noise. In *Proceedings of the IEEE/CVF Winter Conference on applications of computer vision*, pages 3159–3168, 2021.

[60] Thomas C Schelling. *Micromotives and macrobehavior*. WW Norton & Company, 2006.

[61] Hans Ulrich Simon. Bounds on the number of examples needed for learning functions. In John Shawe-Taylor and Martin Anthony, editors, *Proceedings of the First European Conference on Computational Learning Theory, EuroCOLT 1993, London, UK, December 20-22, 1993*, pages 83–94. Oxford University Press, 1993.

[62] Subhrajit Sinha, Sai P Nandanoori, and DA Barajas-Solano. Online real-time learning of dynamical systems from noisy streaming data. *Scientific Reports*, 13(1):22564, 2023.

[63] Michael W Sneddon, James R Faeder, and Thierry Emonet. Efficient modeling, simulation and coarse-graining of biological complexity with nfsim. *Nature methods*, 8(2):177–183, 2011.

[64] S. Soundarajan and J. E. Hopcroft. Recovering social networks from contagion information. In *Proceedings of the 7th Annual Conference on Theory and Models of Computation*, pages 419–430. Springer, 2010.

[65] Daniel Trpevski, Wallace KS Tang, and Ljupco Kocarev. Model for rumor spreading over networks. *Physical Review E*, 81(5):056102, 2010.

[66] Leslie G Valiant. A theory of the learnable. *Communications of the ACM*, 27(11):1134–1142, 1984.

[67] Duncan J Watts. A simple model of global cascades on random networks. *Proceedings of the National Academy of Sciences*, 99(9):5766–5771, 2002.

[68] Zheng Wen, Branislav Kveton, Michal Valko, and Sharan Vaswani. Online influence maximization under independent cascade model with semi-bandit feedback. *Advances in neural information processing systems*, 30, 2017.

[69] Mateusz Wilinski and Andrey Lokhov. Prediction-centric learning of independent cascade dynamics from partial observations. In *International Conference on Machine Learning*, pages 11182–11192. PMLR, 2021.

7 Appendix

In the Appendix, we present details of the technical results established in the main paper. These include (i) proofs for Lemma 3.1, Theorem 3.2, Theorem 4.1, Theorem 4.2, and (ii) pseudocodes for Algorithm 1 (i.e., V-ERM) and Algorithm 3 (i.e., VisRange). We also include additional experimental results for the algorithms.

7.1 Additional information for Section 2

The settings of existing studies on learning dynamical systems

Our setting for learning discrete dynamical systems (see Section 2 in the main paper) aligns with the line of existing research. Here, we present the detailed setting of a few illustrative papers on learning discrete dynamical systems. These works are also cited in our main paper.

All the existing work presented below considers the following setting: 1. The vertex state is **binary**; 2. The update scheme is **synchronous**; 3. The time-scale is **discrete**; 4. The interaction functions are either threshold functions, susceptible-infected, or independent cascade.

1. `AAAI-2022` [13]; 2. `ICML-2022` [55]; 3. `ICML-2021` [10]; 4. `ICML-2021` [16];

5. `ICML-2021` [69]; 6. `NeurIPS-2020` [44]; 7. `ICML-2020` [12]; 8. `ICML-2019` [2];

9. `NeurIPS-2016` [29]; 10. `NeurIPS-2015` [49].

A pictorial example of a discrete dynamical system

We present a toy example of the evolution of a threshold discrete dynamical system. The goal of this figure is to assist readers in understanding the formal definitions presented in Section 2 of the main paper. For large-scale realistic discrete dynamical systems used in the real world, we refer readers to the following references: [7, 33, 46, 63, 60, 42, 35].


Figure 4: The evolution of an example threshold dynamical system with 5 vertices. The threshold value of each vertex is shown in the figure. Here, we present system updates over two time-steps, where vertices in state-1 are highlighted in blue. For instance, the threshold of vertex v_1 is 2. In the first configuration, v_1 has only one neighbor in state-1, which is less than its threshold. Therefore, its state gets updated to 0 (shown in the second configuration) after time step 1. In the second configuration, v_1 has two neighbors of type-1, which satisfies its threshold value. Thus, its state gets updated to 1 after time-step 2.

7.2 Additional information for results in Section 3

The general learning model

We now present the details of the general learning model stated in Section 3. Following this, learning discrete dynamical systems becomes a *special case* of this general scheme. The new scheme follows our learning setting in Section 2 for a *finite* hypothesis class $\mathcal{H}'$, with the following generalizations:

- **Training set.** The training set contains *pairs* of n-dimensional vectors. For each pair, the input vector represents the features of n entities (e.g., vertices); their corresponding labels are computed from the input features *collectively*.

  Formally, a training set $\mathcal{T}$ of size q consists of q pairs of n-dimensional vectors, denoted by $(\mathcal{W}_j, \hat{\mathcal{W}}_j)$, $j = 1, ..., q$. For each $(\mathcal{W}_j, \hat{\mathcal{W}}_j) \in \mathcal{T}$, the input vector $\mathcal{W}_j \in \{0, 1\}^n$ is drawn from an *unknown* distribution $\mathcal{D}$, where each entry of $\mathcal{W}_j$ is the feature value associated with an entity.

  Let $\mathcal{W}'_j$ denote the *true* label for $\mathcal{W}_j$, computed by some unknown ground truth labeling function (i.e., hypothesis) $h^* \in \mathcal{H}'$. Here, each $\mathcal{W}'_j[i] \in \{0, ..., k-1\}$ is the label for the ith entity, $i = 1, ..., n$, where there are $k \geq 2$ possible labels. We do not restrict the actual form of h^*.

- **Noise process**. The observed label vector $\hat{\mathcal{W}}_j$ in $(\mathcal{W}_j, \hat{\mathcal{W}}_j)$ is an erroneous version of $\mathcal{W}'_j$. In particular, for each entity $i = 1, ..., n$ and for some $0 \leq \eta_i < 1/2$, (i) with probability $1 - \eta_i$, the value of $\hat{\mathcal{W}}_j[i] = \mathcal{W}'_j[i]$; (ii) with probability η_i, $\hat{\mathcal{W}}_j[i]$ is a label in $\{0, ..., k-1\} \smallsetminus \{\mathcal{W}'_j[i]\}$ that is *different* from the true label $\mathcal{W}'_j[i]$, chosen uniformly at random from the $k-1$ labels. Let $\bar{\eta} < 1/2$ be an upper bound on the error terms.

The goal is to search for a hypothesis $h \in \mathcal{H}'$ such that, when given a new feature vector $\mathcal{W} \sim \mathcal{D}$, h predicts the resulting true label vector $\mathcal{W}'$ under the PAC guarantee.

Detailed proof of Lemma 3.1.

Recall that Lemma 3.1 establishes an upper bound on the number of noisy training samples needed by any element-wise ERM for learning under the general learning model.

> **Lemma 3.1** *Let $\mathcal{P} = \{\mathcal{P}_1, ..., \mathcal{P}_n\}$ be a collection of partitions of the hypothesis class $\mathcal{H}'$; $t_{max}(\mathcal{P})$ is the size of the largest partition in $\mathcal{P}$. For any $\epsilon, \delta \in (0, 1)$, and any $\eta_i \leq \bar{\eta} < 1/2$, $i = 1, ..., n$, with a training set of size*
> $$q = O\left(\frac{\bar{\eta}}{(1 - 2\bar{\eta})^2} \cdot \frac{1}{\epsilon^2} \cdot n^2 \log(\frac{t_{\max}(\mathcal{P}) \cdot n}{\delta}) \right)$$
> *any element-wise ERM $\mathcal{A}$ learns a hypothesis $h \in \mathcal{H}'$ such that with probability at least $1 - \delta$ (over $\mathcal{T} \sim \mathcal{O}^q$), $err_{\mathcal{D}}(h) < \epsilon$.*

Proof. We show that a training set of size

$$q = \left\lceil \frac{8\bar{\eta} \cdot n^2 + \epsilon(1 - (2 - k') \cdot \bar{\eta}) \cdot n}{\epsilon^2 \cdot (1 - (2 - k') \cdot \bar{\eta})^2} \cdot \ln\left(\frac{2t_{\max}(\mathcal{P}) \cdot n}{\delta} \right) \right\rceil \tag{1}$$

is sufficient to establish the (ϵ, δ)-PAC guarantee, where $k' = (k-2)/(k-1)$.

Let

$$err_{\mathcal{D}}(h,i) = \Pr_{\mathcal{W}\sim\mathcal{D}}[h(\mathcal{W})[i] \neq h^*(\mathcal{W})[i]]$$

be the probability (over $\mathcal{W} \sim \mathcal{D}$) of a hypothesis h makes a wrong prediction on the state of the entity $i = 1, ..., n$. We say that a hypothesis $h \in \mathcal{H}$ is γ-**bad** w.r.t entity i if $err_{\mathcal{D}}(h,i) > \gamma$.

We show the following result.

Claim 7.0.1. The probability (over $\mathcal{T} \sim \mathcal{O}^q$) that the learned hypothesis h is ϵ/n-bad w.r.t. at least one entity $i = 1, .., n$ is at most δ

Note that by the above claim,

$$\Pr_{\mathcal{T}\sim\mathcal{O}^q}[err_{\mathcal{D}}(h) > \epsilon] \leq \delta \tag{2}$$

since if $err_{\mathcal{D}}(h) > \epsilon$ for the learned h, it must be the case that $err_{\mathcal{D}}(h,i) > \epsilon/n$ for at least one entity $i = 1, .., n$. We thus have

$$\Pr_{\mathcal{T}\sim\mathcal{O}^q}[err_{\mathcal{D}}(h) > \epsilon] \leq \Pr_{\mathcal{T}\sim\mathcal{O}^q}[err_{\mathcal{D}}(h,i) > \epsilon \text{ for at least one } i = 1, ..., n] \leq \delta$$

and Lemma 3.1 thus follows. The rest of the proof focuses on showing that Claim 7.0.1 holds under the training size q given in Eq 1.

First, **fix an entity** $i \in \{1, ..., n\}$. For a noisy training data point $(\mathcal{W}, \hat{\mathcal{W}}) \sim \mathcal{O}$ provided by the oracle, let $\hat{err}_{\mathcal{O}}(h,i)$ be the probability of $h(\mathcal{W})$ disagrees with the state of i in $\hat{\mathcal{W}}$ returned from the erroneous oracle $\mathcal{O}$. Formally,

$$\hat{err}_{\mathcal{O}}(h,i) = \Pr_{(\mathcal{W},\hat{\mathcal{W}})\sim\mathcal{O}}[h(\mathcal{W})[i] \neq \hat{\mathcal{W}}[i]]$$

Note that for the ground-truth hypothesis h^*, it holds that $\hat{err}_{\mathcal{O}}(h^*,i) = \eta_i$ since the disagreement only happens when an erroneous state is returned by the oracle $\mathcal{O}$. Similarly, for any hypothesis h that is ϵ/n-bad w.r.t. the entity i, one can easily verify that

$$\hat{err}_{\mathcal{O}}(h,i) = (1 - err_{\mathcal{D}}(h,i)) \cdot \eta_i + err_{\mathcal{D}}(h,i) \cdot (1 - \eta_i) + err_{\mathcal{D}}(h,i) \cdot \eta_i \cdot k' \tag{3}$$
$$\geq \eta_i + \frac{\epsilon}{n} \cdot \left(1 - (2 - k')\bar{\eta}\right)$$

where $k' = (k-2)/(k-1)$.

Let

$$z = \frac{\epsilon}{n}(1 - (2 - k')\bar{\eta})$$

be the lower bound on the difference of the error probability (over $(\mathcal{W}, \hat{\mathcal{W}}) \sim \mathcal{O}$) between h^* and any hypothesis h that is ϵ/n-bad w.r.t. entity i.

We now **fix a hypothesis** h' that is ϵ/n-bad w.r.t. entity i. Note that Claim 7.0.1 trivially holds if no such h' exists. Let $X(h',i,\mathcal{T})$ be the random variable representing the number of disagreements over $\mathcal{T}$ between the states of i predicted by h' and the states of i returned by the erroneous oracle $\mathcal{O}$. That is, $X(h',i,\mathcal{T})$ is number of samples $(\mathcal{W}_j, \hat{\mathcal{W}}_j)$ in $\mathcal{T}$ such that $\hat{\mathcal{W}}_j[i] \neq h(\mathcal{W}_j)[i]$.

Note that if this hypothesis h' is returned by our element-wise ERM $\mathcal{A}$, the empirical loss of h' must be the minimum over all hypotheses in $\mathcal{H}$, including the true hypothesis h^*. Thus, at least one of the following two events must occur: (I) $X(h^*, i, \mathcal{T}) > \eta_i q + z/2 \cdot q$; (II) $X(h', i, \mathcal{T}) \leq \eta_i q + z/2 \cdot q$.

Observe that

$$\mathbb{E}[X(h^*, i, \mathcal{T})] = \eta_i \cdot q, \quad \mathbb{E}[X(h', i, \mathcal{T})] \geq (\eta_i + z)q$$

where the expectation is taken over $\mathcal{T} \sim \mathcal{O}^q$. By Chernoff,

$$\Pr_{\mathcal{T} \sim \mathcal{O}^q}[X(h', i, \mathcal{T}) \leq \eta_i q + z/2 \cdot q] = \Pr_{\mathcal{T} \sim \mathcal{O}^q}[X(h', i, \mathcal{T}) \leq (1 - \frac{z}{2(\eta_i + z)}) \cdot (\eta_i + z)q]$$

$$\leq \exp\left(-\frac{1}{8} \cdot q \cdot \frac{z^2}{\eta_i + z}\right) \tag{4}$$

One can then verify that, for the value of q specified in Eq 1, we have

$$\exp\left(-\frac{1}{8} \cdot q \cdot \frac{z^2}{\eta_i + z}\right) < \frac{1}{2} \cdot \frac{\delta}{t_{\max}(\mathcal{P})n}$$

and thus by Ineq 4,

$$\Pr_{\mathcal{T} \sim \mathcal{O}^q}[X(h', i, \mathcal{T}) \leq \eta_i q + z/2 \cdot q] < \frac{1}{2} \cdot \frac{\delta}{t_{\max}(\mathcal{P})n} \tag{5}$$

Similarly, one can easily verify that for the same q,

$$\Pr_{\mathcal{T} \sim \mathcal{O}^q}[X(h^*, i, \mathcal{T}) > \eta_i q + z/2 \cdot q] < \frac{1}{2} \cdot \frac{\delta}{n} \tag{6}$$

Recall that the hypothesis class $\mathcal{H}'$ admits a collection of partitions $\mathcal{P} = \{\mathcal{P}_1, ..., \mathcal{P}_n\}$. Under each partition $\mathcal{P}_j \in \mathcal{P}$, the empirical loss (under any training set $\mathcal{T}$) w.r.t the entity j for all hypotheses in each subset of $\mathcal{H}'$ are the same. Suppose the event "$X(h', i, \mathcal{T}) \leq \eta_i q + z/2 \cdot q$" occurs for h', then this event *must* also occurred for all other hypotheses that are in the same subset with h' under $\mathcal{P}_i$.

Since the size of the partition $\mathcal{P}_i$ is $t_i \leq t_{max}$, by Ineq 5, it follows that the probability (over $\mathcal{T} \sim \mathcal{O}^q$) of the event "$X(h', v, \mathcal{T}) \leq \eta_i q + z/2 \cdot q$" occurs for at least one ϵ/n-bad hypothesis h' in the space $\mathcal{H}$ is at most $1/2 \cdot \delta/n$.

Combining Ineq 5 and Ineq 6, it follows that the probability (over $\mathcal{T} \sim \mathcal{O}^q$) of either

Event I: "$X(h^*, i, \mathcal{T}) > \eta_i q + z/2 \cdot q$"

or

Event II: "$X(h', i, \mathcal{T}) \leq \eta_i q + z/2 \cdot q$ for at least one ϵ/n-bad hypothesis h'"

occurs is at most δ/n.

Lastly, let h denote the hypothesis returned by the element-wise ERM $\mathcal{A}$. Note that h is ϵ/n-bad (i.e.,

$err_{\mathcal{D}}(h, i) > \epsilon/n$) w.r.t i only if either event **I** or event **II** (or both) happened. Thus, we have

$$\Pr{}_{\mathcal{T} \sim \mathcal{O}^q}[err_{\mathcal{D}}(h, i) > \epsilon/n] \leq \frac{\delta}{n} \tag{7}$$

Subsequently,

$$\Pr{}_{\mathcal{T} \sim \mathcal{O}^q}[err_{\mathcal{D}}(h, i) > \epsilon/n \text{ for at least one entity } i = 1, ..., n] \leq \delta \tag{8}$$

Claim 7.0.1 follows. This concludes the proof. ∎

Pseudocode of Algorithm 1 Here, we present the pseudocode of Algorithm 1 (i.e., V-ERM).

Algorithm 1: `Vertex-wise ERM` (V-ERM)

    Input : A training set $\mathcal{T}$; graph $\mathcal{G}$
    Output A system h from $\mathcal{H}$
    :
1 **for** $v \in \mathcal{V}(\mathcal{G})$ **do**
2   **for** $\tau = 0, 1, ..., \deg_v + 1$ **do**
3         $h_\tau \leftarrow$ a system where the threshold of v is τ
4         $s_\tau \leftarrow \sum_{(\mathcal{C}, \hat{\mathcal{C}}) \in \mathcal{T}} \mathbb{1}\left(h_\tau(\mathcal{C})[v] \neq \hat{\mathcal{C}}[v]\right)$
5   In h, set $\tau_v \leftarrow \arg\min_\tau\{s_\tau\}$ // The threshold of v in h
6 **return** h

Detailed proof of Theorem 3.2

Theorem 3.2 establishes the sufficient number of training samples for Algorithm 1 on learning threshold discrete dynamical systems. This result is an implication of Lemma 3.1. In particular, one can define a partition of the hypothesis class $\mathcal{H}$ as follows. For each vertex $v \in \mathcal{V}$, consider a partition $\mathcal{P}_v = \{\mathcal{H}_0, ..., \mathcal{H}_{\deg_v+1}\}$ of $\mathcal{H}$ such that each hypothesis in $\mathcal{H}_\ell$, $\ell = 1, ..., \deg_v + 1$, has the property: *the threshold of v is ℓ*. One can easily verify that, for any training set $\mathcal{T}$, the empirical loss w.r.t. vertex v for all hypotheses in $\mathcal{H}_\ell$ are the same. Further, note that the size of the partition $t_v = \deg_v + 2$. Subsequently, the value $t_{\max}(\mathcal{P}) = \Delta + 2$ where $\mathcal{P} = \{\mathcal{P}_1, ..., \mathcal{P}_n\}$. By Lemma 3.1, it follows that a training set of size

$$q = \left\lceil \frac{8\bar{\eta} \cdot n^2 + \epsilon(1 - 2\bar{\eta}) \cdot n}{\epsilon^2 \cdot (1 - 2\bar{\eta})^2} \cdot \ln\left(\frac{3\Delta \cdot n}{\delta}\right) \right\rceil$$

is sufficient. Since $\Delta < n$, the theorem follows.

Theorem 3.2 *For any $\epsilon, \delta \in (0, 1)$, and any $\eta_v \leq \bar{\eta} < 1/2, v \in \mathcal{V}$, with a training set of size*

$$q = O\left(\frac{\bar{\eta}}{(1 - 2\bar{\eta})^2} \cdot \frac{1}{\epsilon^2} \cdot n^2 \cdot \log(\frac{n}{\delta})\right)$$

Algorithm 1 (i.e., (i.e., V-ERM)) learns a hypothesis $h \in \mathcal{H}$ such that with probability at least $1 - \delta$ (over $\mathcal{T} \sim \mathcal{O}^q$), $err_{\mathcal{D}}(h) < \epsilon$.

7.3 Additional information for results in Section 4

Pseudocode for Algorithm 2. We present the pseudocode for Algorithm 2 VisScore on learning based on *visiting times of scores*.

Algorithm 2: `Visiting Scores (VisScore)`

    Input : A training set $\mathcal{T}$; graph $\mathcal{G}$
    Output A system h
    :

1 $q \leftarrow |\mathcal{T}|$ `// Size of the training set`
2 **for** $v \in \mathcal{V}(\mathcal{G})$ **do**
3 $\lambda_s \leftarrow 0, s = 0, ..., \deg_v + 1$ `// The hitting time of score`
4 $a_s, b_s \leftarrow 0, s = 0, ..., \deg_v + 1$ `// Count the number of output state-0's and output`
        `state-1's under input score s`
5   **for** $(\mathcal{C}, \hat{\mathcal{C}}) \in \mathcal{T}$ **do**
6 $s \leftarrow \text{score}(\mathcal{C}, v)$
7 $\lambda_s \leftarrow \lambda_s + 1$
8 $a_s \leftarrow a_s + 1$ **if** $\hat{C}[v] == 0$; **else** $b_s \leftarrow b_s + 1$
9 $S \leftarrow \varnothing$
10   **for** $s = 0, ..., \deg_v + 1,$ **do**
11   **if** $\lambda_s \geq \frac{\epsilon}{2\Delta n} \cdot q$ **then**
12 $\ell'_s \leftarrow 0$ **if** $a_s > b_s$; **else** $\ell'_s \leftarrow 1$ `// Majority voting on the correct output`
            `state of v under input score s`
13 $S \leftarrow S \cup \{s\}$
14   **if** $\exists s \in S$ *s.t.* $\ell'_s = 0$ **then**
15   In h, set $\tau_v \leftarrow 1 + \max\{s : s \in S, \ell'_s = 0\}$ `// The learned threshold of v`
16   **else**
17   In h, set $\tau_v \leftarrow 0$
18 **return** h

Detailed proof of Theorem 4.1

In Theorem 4.1, we prove the sufficient number of training samples needed by Algorithm 2.

> **Theorem 4.1** *For any $\epsilon, \delta \in (0, 1)$, and any $\eta_v \leq \bar{\eta} < 1/2, v \in \mathcal{V}$, with a training set of size*
>
> $$q = O\left(\frac{1 - \bar{\eta}}{(1/2 - \bar{\eta})^2} \cdot \frac{1}{\epsilon} \cdot \Delta n \cdot \log(\frac{n}{\delta})\right)$$
>
> *Algorithm 2 (i.e., VisScore) learns a hypothesis $h \in \mathcal{H}$ such that with probability at least $1 - \delta$ (over $\mathcal{T} \sim \mathcal{O}^q$), we have $err_{\mathcal{D}}(h) < \epsilon$.*

Proof. We prove that a training set of size

$$q = 4 \cdot \frac{1 - \bar{\eta}}{(1/2 - \bar{\eta})^2} \cdot \frac{\Delta n}{\epsilon} \cdot \ln\left(\frac{4\Delta n}{\delta}\right) \tag{9}$$

is sufficient for Algorithm 2 to guarantee the (ϵ, δ)-PAC bound. Recall that $\Delta < n$ is the maximum degree of

the underlying graph. From a high level, we show that, when q is sufficiently large, with probability at least $1 - \delta$ over $\mathcal{T} \sim \mathcal{D}^q$, each score with a relatively "high" visiting probability will be visited a sufficiently large number of times. We then take the majority output state of v over the erroneous successors under each such score to be the correct output state and subsequently infer the threshold v.

As shown in Algorithm 2, we learn the set of thresholds in a vertex-wise manner. **Fix a vertex** $v \in \mathcal{V}$. For each of v's possible scores, denoted by $s = 0, ..., \deg_v + 1$, we say that a score s is ϵ-**important** w.r.t. v if its visiting probability (over $\mathcal{C} \sim \mathcal{D}$) is at least ϵ. Recall that the visiting probability of a score s in terms of v is the probability of sampling a configuration $\mathcal{C} \sim \mathcal{D}$ such that the score of v under $\mathcal{C}$ is s.

Now, we **fix a score** s that is $(\epsilon/(\Delta n))$-important w.r.t v. Note that at least one score in $\{0, ..., \deg_v\}$ is $(\epsilon/(\Delta n))$-important w.r.t v. Let $X(s, \mathcal{T})$ be the random variable (over $\mathcal{T} \sim \mathcal{O}^q$) representing the *visiting times* of the score s in the training set $\mathcal{T}$. That is, $X(s, \mathcal{T})$ is the number of $(\mathcal{C}, \hat{\mathcal{C}}) \in \mathcal{T}$ such that the score of v under $\mathcal{C}$ is s.

Importantly, since s is $(\epsilon/(\Delta n))$-important, the expected value of $X(s, \mathcal{T})$ satisfies

$$\mathbb{E}[X(s, \mathcal{T})] \geq \frac{\epsilon}{\Delta n} \cdot q$$

Recall that in Algorithm 2, we only learn from the scores whose visiting time in $\mathcal{T}$ is at least

$$t = \frac{1}{2} \cdot \frac{\epsilon}{\Delta n} \cdot q$$

Ideally, every $(\epsilon/(\Delta n))$-important score is visited at least $t = 1/2 \cdot \epsilon/(\Delta n) \cdot q$ when q is sufficiently large. Specifically, by tail bound, one can verify that

$$\Pr_{\mathcal{T} \sim \mathcal{O}^q}[X(s, \mathcal{T}) < t] \leq \exp\left(-\frac{1}{8} \cdot \frac{\epsilon}{\Delta n} \cdot q\right) \tag{10}$$

where the event "$X(s, \mathcal{T}) < t$" occurring for the $(\epsilon/(\Delta n))$-important score s is undesirable. It then follows that

$$\Pr_{\mathcal{T} \sim \mathcal{Q}^q}[X(s, \mathcal{T}) < t \text{ for at least one } \epsilon/(\Delta n)\text{-important score } s] \tag{11}$$
$$\leq 2\Delta \cdot \exp\left(-\frac{1}{8} \cdot \frac{\epsilon}{\Delta n} \cdot q\right)$$

there there are at most Δ scores for vertex v. For clarity, we define the above (bad) event:

Event I: $X(s, \mathcal{T}) < 1/2 \cdot \frac{\epsilon}{\Delta n} \cdot q$ for at least one $\epsilon/(\Delta n)$-important score s

One can then verify that, when the size of the training set q satisfies Eq 9, Event I happens with probability (over $\mathcal{T} \sim \mathcal{O}^q$) at most $\delta/(2n)$.

Another desirable property that Algorithm 2 utilizes is that, for sufficiently large q, when a score s is visited enough number of times (i.e., at least $t = 1/2 \cdot \epsilon/(\Delta n) \cdot q$ times) in $\mathcal{T}$, the *majority output state* of v over the erroneous successor under the input score s is the *true state* of v in an error-free successor under the

same input score s. We now prove this.

We **fix a score** s of v that got visited at least $t = (1/2) \cdot \epsilon/(\Delta n) \cdot q$ times in $\mathcal{T}$. Note that such a score must exist. Let $Q(s, T)$ be the visiting times for s under $\mathcal{T}$; $Q(s, T) \geq t$. Let ℓ_s be the correct output state of v under the input score s. That is, when there are no errors, if the input to v's interaction function is s, then the state of v returned by the ground-truth system h^* is ℓ_s.

Let $Y(\ell_s, s, \mathcal{T})$ be the number of times that ℓ_s appears as the output state of v over the erroneous successors in the training set $\mathcal{T}$, under the input score s. Note that

$$\mathbb{E}[Y(\ell_s, s, \mathcal{T})] \geq (1 - \bar{\eta}) \cdot Q(s, T) \geq (1 - \bar{\eta}) \cdot t$$

where $\bar{\eta} < 1/2$ is the upper bound on the error terms over all vertices. Ideally, we should have $Y(\ell_s, s, \mathcal{T}) > 1/2 \cdot Q(s, T)$, that is, the state ℓ_s wins is the majority vote. By Chernoff, one can verify that:

$$\Pr_{\mathcal{T} \sim \mathcal{O}^q}[Y(\ell_s, s, \mathcal{T}) \leq \frac{1}{2} \cdot Q(s, T)] \tag{12}$$
$$= \Pr_{\mathcal{T} \sim \mathcal{O}^q}[Y(\ell_s, s, \mathcal{T}) \leq (1 - \frac{(1 - \bar{\eta}) - 1/2}{1 - \bar{\eta}}) \cdot ((1 - \bar{\eta}) \cdot Q(s, T))]$$
$$\leq \exp\left(-\frac{1}{2} \cdot \left(\frac{(1 - \bar{\eta}) - 1/2}{1 - \bar{\eta}}\right)^2 \cdot (1 - \bar{\eta}) \cdot Q(s, T)\right)$$
$$\leq \exp\left(-\frac{1}{2} \cdot \left(\frac{(1 - \bar{\eta}) - 1/2}{1 - \bar{\eta}}\right)^2 \cdot (1 - \bar{\eta}) \cdot \frac{\epsilon q}{2 \Delta n}\right)$$

where the event "$Y(\ell_s, s, \mathcal{T}) \leq \frac{1}{2} \cdot Q(s, T)$" occurring for the label ℓ_s is undesirable. It follows that

$$\Pr_{\mathcal{T} \sim \mathcal{O}^q}[Y(\ell_s, s, \mathcal{T}) \leq \frac{1}{2} \cdot Q(s, T) \text{ for at least one score } s \text{ with visiting time at least } t] \tag{13}$$
$$\leq 2\Delta \cdot \exp\left(-\frac{1}{2} \cdot \left(\frac{(1 - \bar{\eta}) - 1/2}{1 - \bar{\eta}}\right)^2 \cdot (1 - \bar{\eta}) \cdot \left(\frac{1}{2} \frac{\epsilon}{\Delta n} q\right)\right)$$

Let Event II be the above bad event:

Event II: $Y(\ell_s, s, \mathcal{T}) \leq 1/2 \cdot Q(s, T)$ for at least one score s with visiting time at least t

One can verify that for a training set of size q shown in Eq 9, we have that Event II happens with probability (over $\mathcal{T} \sim \mathcal{O}^q$) at most $\delta/(2n)$.

Let h be the hypothesis returned by Algorithm 2. We say that hypothesis h is ϵ/n-**good** w.r.t v if $\Pr_{\mathcal{C} \sim \mathcal{D}}[h(\mathcal{C})[v] \neq h^*(\mathcal{C})[v]] < \epsilon/n$. We remark that:

Claim 7.0.2. If **both** Event I and Event II do **not** occur under a training set $\mathcal{T}$, then the learned h must be ϵ/n-good w.r.t v.

To see this, suppose that both Event I and Event II do **not** occur. Then, the following are true simultaneously:

(i) Every $\epsilon/(\Delta n)$-important score got visited at least t times.

(ii) For any scores that got visited at least t times, the majority voting scheme (over $\mathcal{T}$) gives the correct output state of v under the input score s to v's interaction function.

Let S be the set of scores that got visited at least t times. Note that S includes all the $(\epsilon/(\Delta n))$-important scores, plus possibly some other scores.

Since both Event I and Event II do **not** occur, h learns the correct output state of v under each input score in S. Let $S' \subseteq S$ be the subset of scores such that the learned output state of v is 0. If S' is an empty set, we then have $\tau_v = 0$ as stated in Algorithm 2. On the other hand, suppose S' contains at least one score. As shown in Algorithm 2, we then set

$$\tau_v = \max{}_{s \in S'}\{s\} + 1$$

For either case, one can easily verify that h would **not** make a wrong prediction on the output state of v when seeing any scores in S as an input. Consequently,

*For a $\mathcal{C} \sim \mathcal{D}$, the hypothesis h can make a **wrong** prediction on the output state of v only when the input score (under $\mathcal{C}$) is **not** in S.*

Note that any score not in S is **not** $(\epsilon/(\Delta n))$-important w.r.t. v. That is, the probability of visiting a score $s \notin S$ of v under $\mathcal{C} \sim \mathcal{D}$ is less that $\epsilon/(\Delta n)$. Since there are at most Δ such "bad" scores, we have $\Pr_{\mathcal{C} \sim \mathcal{D}}[h(\mathcal{C})[v] \neq h^\star(\mathcal{C})[v]] < \epsilon/n$. This concludes the claim.

Lastly, since Event I (and Event II) happen with probability (over $\mathcal{T} \sim \mathcal{O}^q$) at most $\delta/(2n)$ w.r.t v, the probability of either one of them happening is at most δ/n. It follows that, the probability that h is (ϵ/n)-bad w.r.t v (i.e., $\Pr_{\mathcal{C} \sim \mathcal{D}}[h(\mathcal{C})[v] \neq h^\star(\mathcal{C})[v]] < \epsilon/n$) is at most δ/n. That is,

$$\Pr{}_{\mathcal{T} \sim \mathcal{O}^q}[h \text{ is } (\epsilon/n)\text{-bad w.r.t at least one } v \in \mathcal{V}] \leq \delta \tag{14}$$

Thus, the probability (over $\mathcal{T} \sim \mathcal{O}^q$) of $err_{\mathcal{D}}(h) > \epsilon$ is at most δ. This concludes the proof. ∎

Pseudocode for Algorithm 3. We present the pseudocode for Algorithm 3 VisRange on learning based on *visiting times of ranges of scores*.

Algorithm 3: `Visiting Ranges (VisRange)`

Input : A training set $\mathcal{T}$; graph $\mathcal{G}$

Output A system h

:

1   $q \leftarrow |\mathcal{T}|$ `// Size of the training set`

2   **for** $v \in \mathcal{V}(\mathcal{G})$ **do**

3      $\lambda_{s_1,s_2} \leftarrow 0$, for $s_1, s_2 = 0, ..., \deg_v + 1, s_1 \leq s_2$ `// The hitting time of range`

4      $a_{s_1,s_2}, b_{s_1,s_2} \leftarrow 0$, for $s_1, s_2 = 0, ..., \deg_v + 1, s_1 \leq s_2$ `// The number of output state-0's`
        `and output state-1's under input scores in range` $[s_1, s_2]$

5   **for** $(\mathcal{C}, \hat{\mathcal{C}}) \in \mathcal{T}$ **do**

6         $s \leftarrow \texttt{score}(\mathcal{C}, v)$

7         $\lambda_{s_1,s_2} \leftarrow \lambda_{s_1,s_2} + 1$ for each range $[s_1, s_2]$ that contains s

8   **For** each range $[s_1, s_2]$ that contains s: $a_{s_1,s_2} \leftarrow a_{s_1,s_2} + 1$ **if** $\hat{C}[v] == 0$; **else**
         $b_{s_1,s_2} \leftarrow b_{s_1,s_2} + 1$

9      $S \leftarrow \varnothing$

10   **for** $s_1 = 0, ..., \deg_v + 1$ **do**

11   **for** $s_2 = s_1, ..., deg_v + 1$ **do**

12   **if** $\lambda_{s_1,s_2} \geq \frac{\epsilon}{2n} \cdot q$ **then**

13              $\ell'_{s_1,s_2} \leftarrow 0$ if $a_{s_1,s_2} > b_{s_1,s_2}$; Else, $\ell'_{s_1,s_2} \leftarrow 1$ `// Majority voting on the`
               `correct output state of` v `under input scores in range` $[s_1, s_2]$

14           $S \leftarrow S \cup \{(s_1, s_2)\}$

15   **if** $\exists (s_1, s_2) \in S$ s.t. $\ell'_{s_1,s_2} = 0$ **then**

16   In h, set $\tau_v \leftarrow 1 + \max\{s_1 : (s_1, s_2) \in S, \ell'_{s_1,s_2} = 0\}$ `// The learned threshold`

17   **else**

18   In h, set $\tau_v \leftarrow 0$

19   **return** h

Detailed proof of Theorem 4.2

In Theorem 4.2, we prove the sufficient number of training samples needed by Algorithm 3.

> **Theorem 4.2** *For any $\epsilon, \delta \in (0, 1)$, and any $\eta_v \leq \bar{\eta} < 1/2, v \in \mathcal{V}$, with a training set of size*
>
> $$q = O\left(\frac{1 - \bar{\eta}}{(1/2 - \bar{\eta})^2} \cdot \frac{1}{\epsilon} \cdot n \cdot \log(\frac{n}{\delta}) \right)$$
>
> *Algorithm 3 (i.e., VisRange) learns a hypothesis $h \in \mathcal{H}$ such that with probability at least $1 - \delta$ (over $\mathcal{T} \sim \mathcal{O}^q$), we have $err_{\mathcal{D}}(h) < \epsilon$.*

Proof. We show that a training set of size

$$q = 8 \cdot \frac{1 - \bar{\eta}}{(1/2 - \bar{\eta})^2} \cdot \frac{n}{\epsilon} \cdot \ln\left(\frac{2\Delta n}{\delta}\right) \tag{15}$$

is sufficiently large to establish the (ϵ, δ)-PAC guarantee. Recall that $\Delta < n$ is the maximum degree of the underlying graph. Conceptually, different from the analysis of Algorithm 2 where one cares about the number of times each *score* is visited in a training set $\mathcal{T}$, here in Algorithm 3, we focus on the visiting time of each

possible *range of scores*. In particular, for a vertex v, each range R of scores with a "high" visiting probability should be visited a large number of times. Consequently, the majority vote over the output states of v in the successors under each score in R would reveal information about the true threshold of v, which is captured by Algorithm 3.

Fix a vertex $v \in \mathcal{V}$. We consider all possible ranges of scores of v, denoted by

$$R_{s_1,s_2} = [s_1, s_2], s_1, s_2 = 0, ..., \deg_v + 1, s_1 \le s_2$$

We remark that there are $O(\Delta^2)$ such ranges for v, where Δ is the maximum degree of the graph.

Similar to the definition of "importance" for scores, we say that a range R_{s_1,s_2} is ϵ-**important** w.r.t. v, if the visiting probability (over $\mathcal{C} \sim \mathcal{D}$) of R_{s_1,s_2} is at least ϵ. Recall that the visiting probability of a range R_{s_1,s_2} w.r.t. v is the probability of sampling a configuration $\mathcal{C} \sim \mathcal{D}$ such that the score of v under $\mathcal{C}$ is in R_{s_1,s_2}.

Fix a range R_{s_1,s_2} that is (ϵ/n)-important w.r.t v; at least one such a range R_{s_1,s_2} exists. Let $X(R_{s_1,s_2}, \mathcal{T})$ be the random variable (over $\mathcal{T} \sim \mathcal{O}^q$) representing *the sum of the visiting times* over the scores in the range R_{s_1,s_2} under training set $\mathcal{T}$. That is,

$$X(R_{s_1,s_2,}, \mathcal{T}) = \sum_{s=s_1}^{s_2} X(s, \mathcal{T}) \tag{16}$$

where $X(s, \mathcal{T})$ is the random variable (over $\mathcal{T} \sim \mathcal{O}^q$) that records the visiting time of each score s under training set $\mathcal{T}$. Given that R_{s_1,s_2} is (ϵ/n)-important, the expected value of $X(R_{s_1,s_2,}, \mathcal{T})$ satisfies that

$$\mathbb{E}[X(R_{s_1,s_2,}, \mathcal{T})] \ge \frac{\epsilon}{n} \cdot q$$

In Algorithm 3, one only cares about the ranges whose visiting time under $\mathcal{T}$ is at least

$$t = \frac{1}{2} \cdot \frac{\epsilon}{n} \cdot q$$

When q is sufficiently large, we want every (ϵ/n)-important range to be visited at least t times so these important ranges will be examined by the learning algorithms. By Chernoff, one can verify that

$$\Pr_{\mathcal{T} \sim \mathcal{O}^q}[X(R_{s_1,s_2,}, \mathcal{T}) < t] \le \exp\left(-\frac{1}{8} \cdot \frac{\epsilon}{n} \cdot q\right) \tag{17}$$

Note that "$X(R_{s_1,s_2,}, \mathcal{T}) < t$" happening for the ϵ/n-important range R_{s_1,s_2} is undesirable. We now define the following bad event:

Event I: $X(R_{s_1,s_2,}, \mathcal{T}) < t$ for at least one ϵ/n-important range R_{s_1,s_2}

By Ineq 17, we have

$$\Pr_{\mathcal{T}\sim\mathcal{Q}^q}\left[X(R_{s_1,s_2},\mathcal{T}) < t \text{ for at least one } \epsilon/n\text{-important range } R_{s_1,s_2}\right] \tag{18}$$

$$\leq 2\Delta^2 \cdot \exp\left(-\frac{1}{8}\cdot\frac{\epsilon}{n}\cdot q\right)$$

where the factor Δ^2 comes from the fact that there are $O(\Delta^2)$ such ranges for v. Importantly, one can verify that, when the size of the training set q satisfies Eq 15, Event I happens with probability (over $\mathcal{T}\sim\mathcal{O}^q$) at most $\delta/(2n)$.

The second property that Algorithm 3 uses is that, when a range R_{s_1,s_2} is visited a sufficiently large number of times (i.e., at least $t = 1/2\cdot\epsilon/n\cdot q$ times) over the training set $\mathcal{T}$, if the *majority output state* of v over all the erroneous successors under the scores in R_{s_1,s_2} is 0, then there must exist at least one score s in R_{s_1,s_2} such that the true output state of v in an error-free successor under the score s is also 0.

We now **fix a range** R_{s_1,s_2} that satisfies the following two properties under a training set $\mathcal{T}\sim\mathcal{O}^q$:

Property I: Both $s_1, s_2 < \tau_v^*$ or both $s_1, s_2 \geq \tau_v^*$, where τ_v^* is the threshold of vertex v under the ground-truth system h^*

Property II: Range R_{s_1,s_2} got visited at least $t = (1/2)\cdot(\epsilon/n)\cdot q$ times.

We remark that such a range R_{s_1,s_2} must exist since the entire range $R_{0,\deg_v+1}$ is visited exactly q times over $\mathcal{T}$. Let $Q(R_{s_1,s_2},\mathcal{T})$ be the visiting times of the range R_{s_1,s_2}; $Q(R_{s_1,s_2},T) \geq t$.

Note that, by property (I) stated above, the *true* output state of vertex v under all the scores in R_{s_1,s_2} are the same. Let ℓ_{s_1,s_2} denote this true output state of vertex v under the scores in R_{s_1,s_2}. That is, when there are no errors, if the input to v's interaction function is any score s in R_{s_1,s_2}, then the state of v returned by the ground-truth system h^* is ℓ_{s_1,s_2}.

Let $Y(\ell_{s_1,s_2},\mathcal{T})$ be the number of training samples $(\mathcal{C},\hat{\mathcal{C}})\in\mathcal{T}$ such that, the score of v under $\mathcal{C}$ is in range $\mathcal{R}_{s_1,s_2}$, and $\hat{C}[v] = \ell_{s_1,s_2}$ is the succeeding state of v. One can verify that

$$\mathbb{E}[Y(\ell_{s_1,s_2},\mathcal{T})] \geq (1-\bar{\eta})\cdot Q(R_{s_1,s_2},\mathcal{T}) \geq (1-\bar{\eta})\cdot t$$

where $\bar{\eta} < 1/2$ is the upper bound on the error terms. Ideally, $Y(\ell_{s_1,s_2},\mathcal{T})$ should be strictly larger than $1/2\cdot Q(R_{s_1,s_2},\mathcal{T})$. That is, the output state ℓ_{s_1,s_2} of v appears in strictly more than half of the pairs $(\mathcal{C},\hat{\mathcal{C}})\in\mathcal{T}$ where $\texttt{score}(\mathcal{C},v)\in R_{s_1,s_2}$. By Chernoff, one can verify that:

$$\Pr_{\mathcal{T}\sim\mathcal{O}^q}\left[Y(\ell_{s_1,s_2},\mathcal{T}) \leq \frac{1}{2}\cdot Q(R_{s_1,s_2},\mathcal{T})\right] \tag{19}$$

$$\leq \exp\left(-\frac{1}{2}\cdot\left(\frac{(1-\bar{\eta})-1/2}{1-\bar{\eta}}\right)^2\cdot(1-\bar{\eta})\cdot\frac{\epsilon q}{2n}\right) \tag{20}$$

Here, the event "$Y(\ell_{s_1,s_2},\mathcal{T}) \leq \frac{1}{2}\cdot Q(R_{s_1,s_2},\mathcal{T})$" happening for label ℓ_{s_1,s_2} is undesirable. Now define the following bad event:

Event II: $Y(\ell_{s_1,s_2},\mathcal{T}) \leq 1/2\cdot Q(R_{s_1,s_2},\mathcal{T})$ for at least one range R_{s_1,s_2} with Property I and II.

By Ineq 19, we have:

$$\Pr{}_{\mathcal{T}\sim\mathcal{O}^q}[Y(\ell_{s_1,s_2},\mathcal{T}) \le 1/2 \cdot Q(R_{s_1,s_2},\mathcal{T}) \text{ for at least one range } R_{s_1,s_2} \text{ with Property I and II}] \quad (21)$$

$$\le 2\Delta^2 \cdot \exp\left(-\frac{1}{2} \cdot \left(\frac{(1-\bar{\eta})-1/2}{1-\bar{\eta}}\right)^2 \cdot (1-\bar{\eta}) \cdot \frac{\epsilon q}{2n}\right)$$

Subsequently, one can verify that for q in Eq 15, Event II happens with probability (over $\mathcal{T} \sim \mathcal{O}^q$) at most $\delta/(2n)$.

Let h be the hypothesis returned by Algorithm 3. We now present the last piece of the proof which shows that h is ϵ/n-good w.r.t v with probability at least $1-\delta$. Recall that h is ϵ/n-good w.r.t v if $\Pr_{\mathcal{C}\sim\mathcal{D}}[h(\mathcal{C})[v] \ne h^*(\mathcal{C})[v]] < \epsilon/n$. The key claim is as follows:

Claim 7.0.3. If **both** Event I and II do **not** occur under $\mathcal{T}$, then the learned h must be ϵ/n-good w.r.t v.

Suppose Claim 7.0.3 is true. We have shown that, under q in Eq 15, Event I (and also Event II) happens with probability (over $\mathcal{T} \sim \mathcal{O}^q$) at most $\delta/(2n)$ w.r.t v. Thus, the probability of either one of them happening is at most δ/n. Consequently, $\Pr_{\mathcal{T}\sim\mathcal{O}^q}[h \text{ is } (\epsilon/n)\text{-bad w.r.t at least one } v] \le \delta$. It follows that, the probability of $err_{\mathcal{D}}(h) > \epsilon$ is at most δ and the proof of the Theorem is complete.

We now show that Claim 7.0.3 is true. If **both** Event I and Event II do **not** occur under $\mathcal{T}$, then the followings are true:

Fact 1. Every ϵ/n-important range is visited at least t times in $\mathcal{T}$.

Let R be the subset consisting of each range R_{s_1,s_2} where (i) both s_1 and s_2 lie on the same side of τ_v^*, and (ii) R_{s_1,s_2} got visited at least t times (note that R_{s_1,s_2} does **not** need to be ϵ/n-important). Note that the true output state of vertex v under all the scores in R_{s_1,s_2} are the same.

Fact 2. For each $R_{s_1,s_2} \in R$, the majority output state of v is the true output state of v.

Let R_0 and R_1 be the subsets of R where the majority output state of v for every range in R_i is i, $i = 0, 1$. Let W be the set of ranges (of scores for v) that got visited at least t times in $\mathcal{T}$, and the corresponding majority output state of v is 0. We remark that $R_0 \subseteq W$. In Algorithm 3, we effectively choose the range in W with the largest s_1 value, denoted by $R_{s_1^*,s_2^*}$ and learn the threshold of v to be $s_1^* + 1$.

Importantly, we observe that such a $R_{s_1^*,s_2^*}$ is **not** in R_1 by Fact 2 above. Thus, it holds that

$$R_{s_1^*,s_2^*} \in R_0, \text{ or } s_1^* < \tau_v^* \le s_2^*$$

Let τ_v' be an positive integer less than τ_v^* where (i) the probability (over $\mathcal{C} \sim \mathcal{D}$) of sampling a configuration $\mathcal{C}$ with score of v in range $[\tau_v', \tau_v^* - 1]$ is larger than ϵ/n, and if $\tau_v' \le \tau_v^* - 2$, then (ii) the probability of sampling a score between $[\tau_v' + 1, \tau_v^* - 1]$ is at most ϵ/n. If no such a τ' exists, then the algorithm sets $\tau_v = 0$ and one can easily verify that the learned h is ϵ/n-good w.r.t v.

Note that the range $R_{\tau'_v, \tau^*_v - 1}$ is an ϵ/n-important range w.r.t v, and both $\tau'_v, \tau^*_v - 1$ lie on the same (left) size of τ^*_v. Thus by Fact 1 and the definition of R_0, we have that

$$R_{\tau'_v, \tau^*_v - 1} \in R_0$$

Since $R_{s^*_1, s^*_2}$ is the range in W with the largest s_1 value, and since $R_0 \subseteq W$ it follows that $s^*_1 \geq \tau'_v$, and the learned threshold

$$\tau_v = s^*_1 + 1 > \tau'_v$$

As a result, the learned system h can only make a wrong prediction on the output state of v when the score of v under the sample configuration $\mathcal{C} \sim \mathcal{D}$ falls within the range $[s^*_1 + 1, \tau_v]$, which happens with probability strictly less than ϵ/n. That is, h is ϵ/n-good w.r.t v. This concludes the proof of Claim 7.0.3 and therefore the theorem. ∎

7.4   Additional details about the experiments

System Specification. All experiments were conducted on an HPC cluster. Each compute node in a cluster is a 20–40 core Intel(R) Xeon(R) CPU E5-2630 v3 @ 2.40GHz processor with 128–384 GB of memory. To achieve scalability, experiments with different networks were conducted on different compute nodes utilizing the SLURM scheduler. Each job used up to 64GB of memory and 1 CPU core during execution. All scripts were implemented using Python 3.7.

Parameter settings of the synthetic networks.

Table 1: Parameter Settings of the Synthetic Networks

Parameter	Notation	Parameter Space
Network Size	n	$\{500, 1000, 2000, 4000, 5000\}$
Average Degree	d_{avg}	$\{5, 10, 20, 30, 40\}$
Noise	η	$\{0.05, 0.1, 0.2, 0.3, 0.4\}$

Information about the networks.

Synthetic Networks. We generated $G_{n,p}$ random graphs based on the ER model [22]. The values of n are shown in Table 1. By using suitable probability p, we also created networks with different average degree (d_{avg}) values, as shown in Table 1. Thus, our experiments use both sparse and dense synthetic networks.

Real-World Networks. The first real-world network we use, *ca-GrQc*, is a collaboration network within the field of General Relativity and Quantum Cosmology spanning published works from January 1993 to April 2003 [43]. It has $5,242$ vertices and $14,496$ edges. The second real-world network, *USpowerGrid*, has $4,941$ vertices and $6,594$ edges. It contains information on the power grid of the Western States of the USA. Here, vertices represent electrical components (e.g., transformer, generator) and edges represent power supply lines [40]. Both of these graphs were obtained from the website (https://net.science) for the net.science software tool [4].

Execution Times. Figure 5 shows the execution times of the algorithms over networks of different sizes

and densities. We ran each algorithm with 100 training samples to maintain consistency in the study of execution times. In general, V-ERM has the longest execution time, followed by VisRange and VisScore. This is consistent with the time complexities of the algorithms. Algorithm 1 (V-ERM) has a time complexity of $\mathcal{O}(n\Delta q)$, where n is the network size, Δ is the maximum degree and q is the size of the training set. On the other hand, (VisScore) Algorithm 2 has a time complexity of $\mathcal{O}(nq + n\Delta)$ and Algorithm 3 (VisRange) has a time complexity of $\mathcal{O}(nq + n\Delta^2)$. For sparse networks (where Δ is a constant), both of these algorithms have a time complexity of $\mathcal{O}(nq)$, as reflected in Figure 5. However, as the network becomes denser, the difference in the second component of the complexity becomes more dominant and this is also reflected in the second panel of Figure 5.
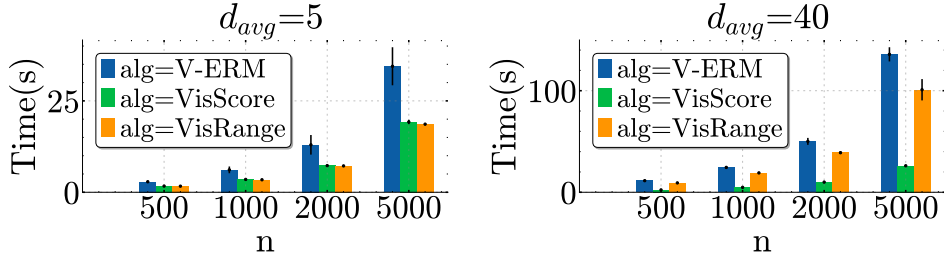

Figure 5: Execution times of the algorithms. The left panel is for sparse networks of various sizes. The right panel is for dense networks. The y-axis scales are different across the two plots. Error bars account for two standard deviation error.


Figure 6: Comparison between the three algorithms on sparse networks in low-noise setting

Comparison of the algorithms: Figure 6 compares the learning patterns of all three algorithms. Contrary to the steady nature of V-ERM, both VisScore and VisRange do not learn reasonably until they have a sufficient number of samples, after which they both learn in a phase-transition like manner. It is also observed that for $N = 2,000$, VisRange requires fewer training samples to achieve reasonable performance compared to VisScore.

Figure 7 shows the number of samples required to achieve threshold loss for the two real-world networks. Since the density and network size are fixed, we only present the sensitivity to noise in the figure. Similar to synthetic networks, we find that the required number of samples increases as the threshold loss is decreased. Moreover, scenarios with higher noise need more training samples for achieving the same threshold loss. It is also observed that the required number of samples for *ca-GrQc* is higher than *UspowerGrid* for same parameter setting. Since *ca-GrQc* is larger than *USpowerGrid* both in terms of network size and density, this behavior coincides with that on synthetic networks, described in the main section of the paper.

V-ERM on synthetic networks: Figure 8 shows the empirical loss of V-ERM for various synthetic networks.
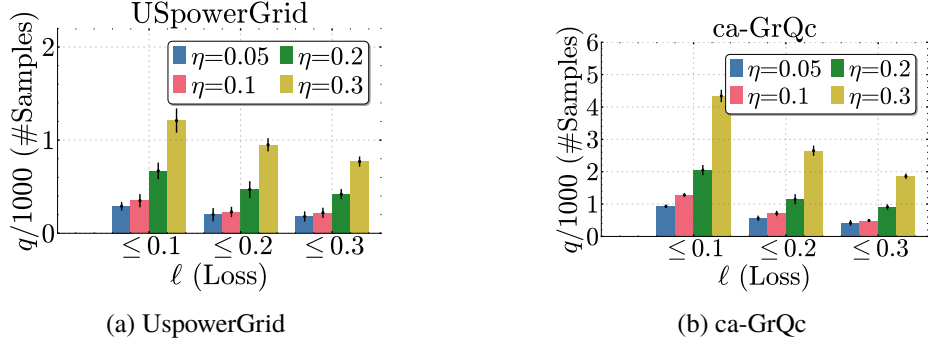
(a) UspowerGrid

(b) ca-GrQc

Figure 7: Training samples needed by V-ERM for achieving specified loss thresholds on real-world networks under various noise settings. Error bars account for two standard deviation errors.

The network size changes along columns and one row of the figure shows networks of same size with increasing density. In general, the learning sensitivity to noise is more prominent for dense networks than sparse networks. This behavior is present but less prominent for different network sizes with same density.
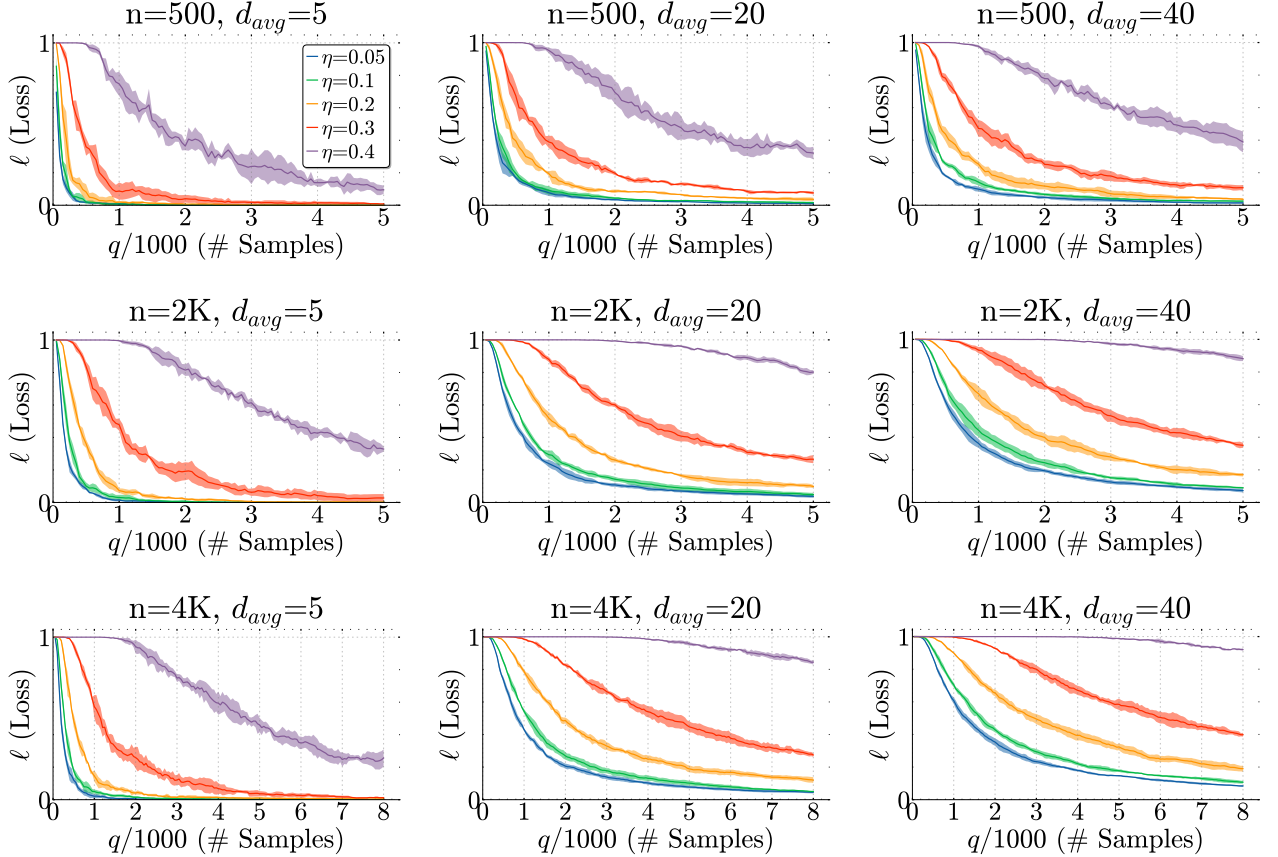

Figure 8: Empirical Evaluation of the learning process of V-ERM (Algorithm 1) under different noise settings for networks of various sizes and densities. Here, $n \in \{500, 2000, 4000\}$ and $d_{avg} \in \{5, 20, 40\}$. The shaded region accounts for 1 standard deviation of the empirical loss.